



Overview of Rolls

Rocks-A-Palooza I

Track 2

Session 1

Rocks Philosophy

- ◆ We've developed a "cluster compiler"
 - ⇒ XML framework + XML parser + kickstart file generator
 - ⇒ Source code + preprocessor + linker
- ◆ Think about "programming your cluster"
 - ⇒ Not "administering your cluster"

Rocks Philosophy

- ◆ Rocks is the “syringe” for software deployment
- ◆ You determine the function of the cluster by selecting Rolls
- ◆ Then you “push” the deployment to all nodes



Goal of Rolls

- ◆ Develop a method to reliably install software on a frontend
- ◆ “User-customizable” frontends
- ◆ Two established approaches:
 - ➔ Add-on method
 - ➔ Rocks method

Add-on Method

1. User responsible for installing and configuring base software stack on a frontend
2. After the frontend installation, the user downloads 'add-on' packages
3. User installs and configures add-on packages
4. User installs compute nodes

Major issue with add-on method

- ◆ The state of the frontend before the add-on packages are added/configured is **unknown**

Rocks Method

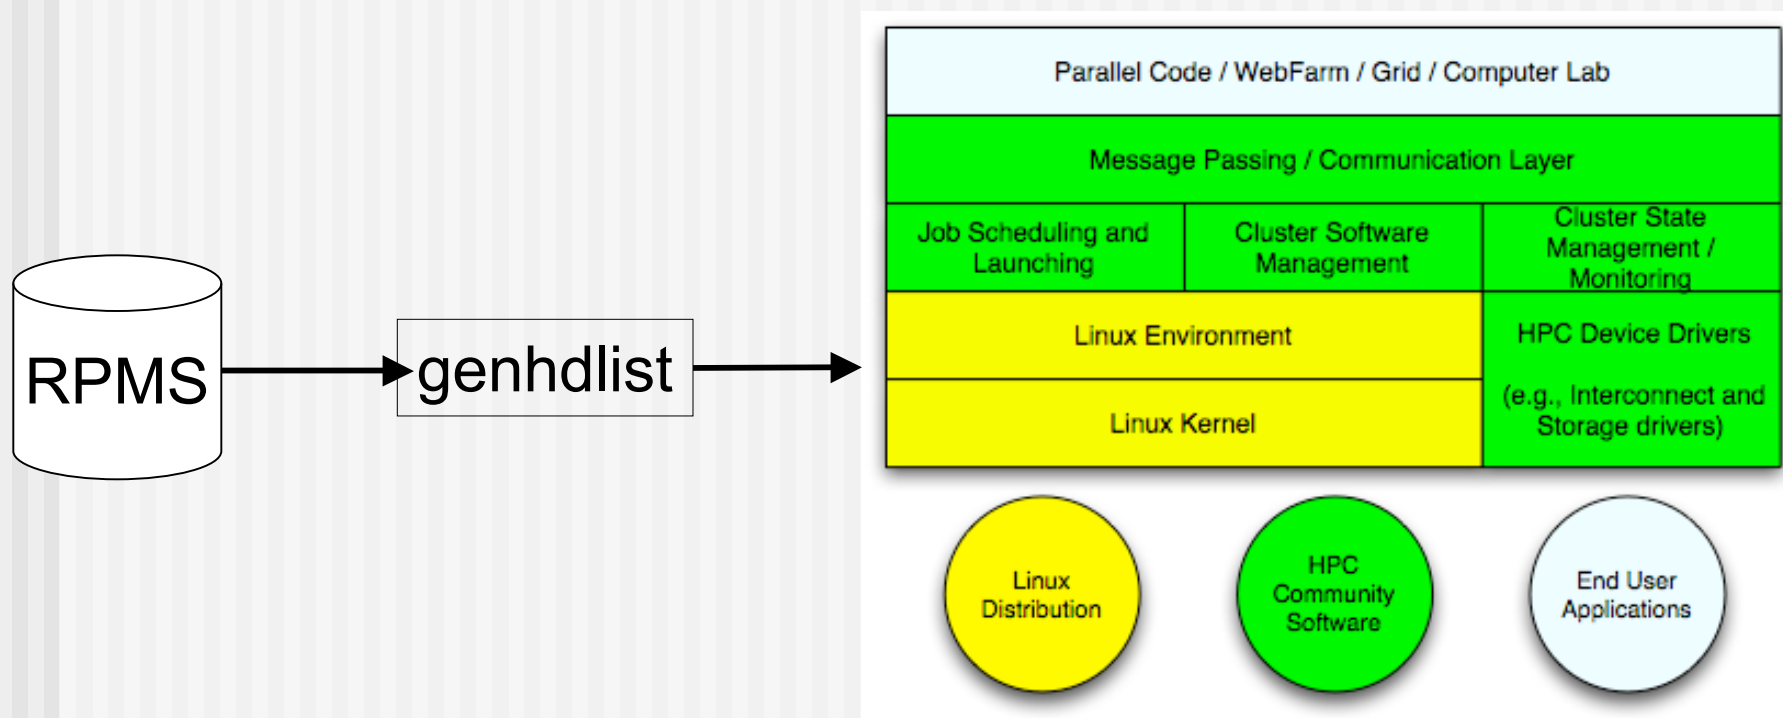
- ◆ To address the major problem with the add-on method, we had the following idea:
 - ➔ All non-RedHat packages must be installed and configured in a **controlled environment**
- ◆ A controlled environment has a known state
- ◆ We chose the RedHat installation environment for the controlled environment

Goal of Rolls

- ◆ This led to modifying the standard RedHat installer in order to accept new packages and configuration
- ◆ A tricky proposition
 - A RedHat distribution is a **monolithic** entity
 - It's tightly-coupled
 - A program called “genhdlist” creates a binary files (hdlist and hdlist2) that contain metadata about every RPM in the distribution
- ◆ To add/remove/change an RPM, you need to re-run genhdlist
 - Else, the RedHat install will not recognize the package
 - Or worse, it fails during package installation



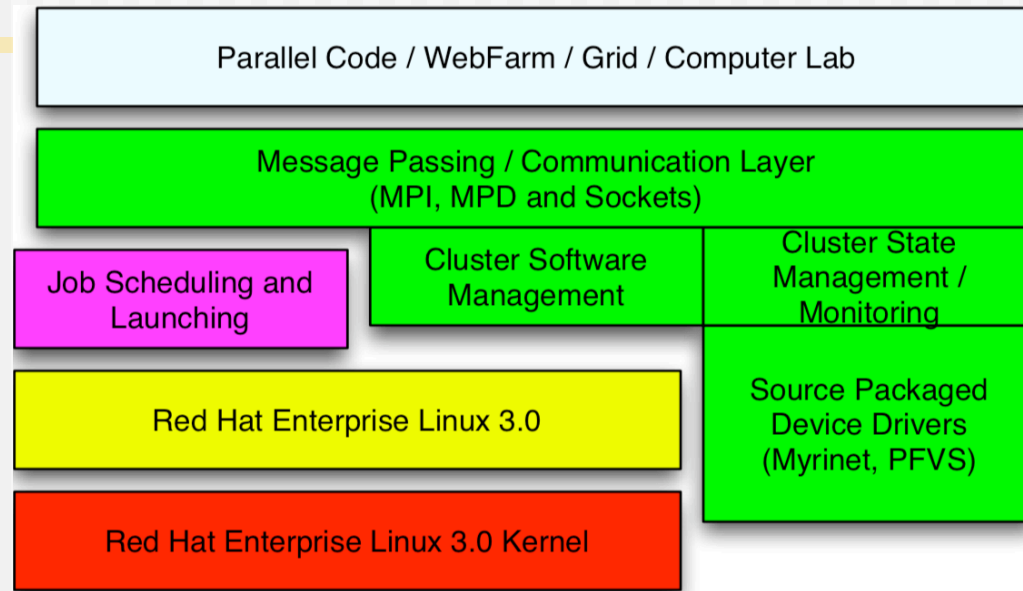
Monolithic Software Stack



Goal of Rolls

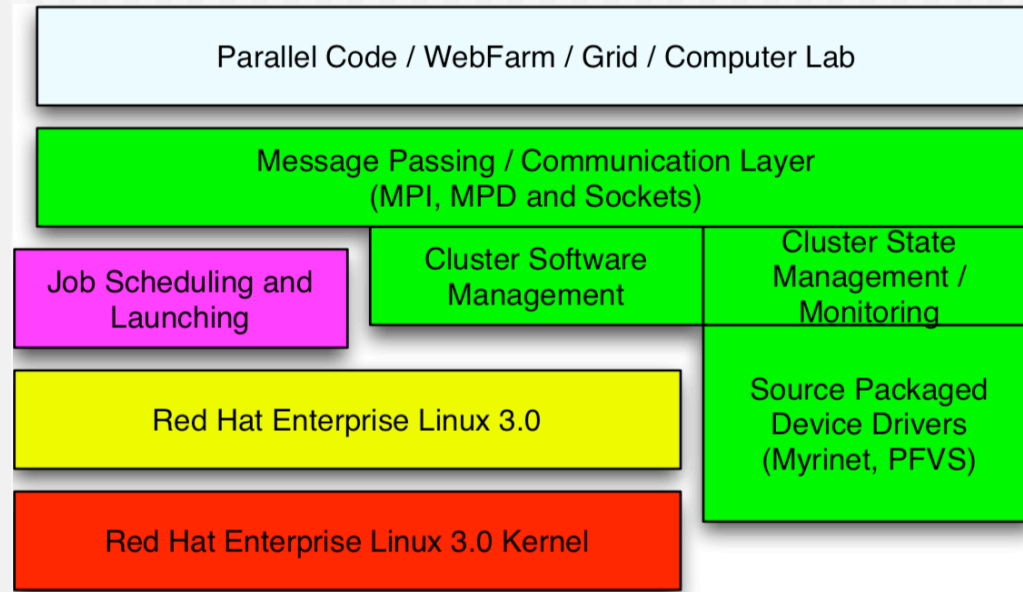
- ◆ Problem: To make the frontend user-customizable at installation time, we needed a mechanism that could accept new packages
- ◆ And, we still wanted to leverage the RedHat installer
 - ➔ We don't want to be in the installer business
- ◆ Solution: Our implementation makes the RedHat installer “think” it is just installing a monolithic RedHat distribution

Goal of Rolls



- ◆ How do you make all the packages above look like a monolithic distribution?
 - Easy! Just run “genhdlist” at release time!
- ◆ But, how do you do it when some of the above blocks are optional and/or unknown?
 - An “unknown” block is one produced after the release or by a third-party

Rolls Function and Value



- ◆ Function: Rolls extend/modify stock RedHat
- ◆ Value: Third parties can extend/modify Rocks
 - ➔ Because Rolls can be optional



The RedHat Installer

Anaconda: RedHat's Installer

- ◆ Open-source python-based installer
- ◆ Developed by RedHat
- ◆ (Somewhat) object-oriented
 - ⇒ We extend when we can and insert “shims” when we can't

Anaconda: RedHat's Installer

◆ Key tasks:

- ⇒ Probe hardware
- ⇒ Ask users for site-specific values
 - E.g., IP addresses and passwords
- ⇒ Insert network and storage drivers
 - For network-based installations and to write packages down onto local disk
- ⇒ Install packages
 - RPMs
- ⇒ Configure services
 - Via shell scripts

Scripted Installation

- ◆ Anaconda achieves “lights-out” installation via **kickstart** mechanism
- ◆ It reads a “kickstart file”
 - ⇒ Description of how to install a node
- ◆ One file composed of three key sections:
 - ⇒ Main: general parameters
 - ⇒ Packages: list of RPMs to install
 - ⇒ Post: scripts to configure services



Kickstart File

◆ Main section

```
rootpw --iscrypted oijewfl5853812kfd
url --url http://10.1.1.1/install/rocks-dist/lan/enterprise/3/en/os/i386
zerombr yes
bootloader --location=mbr
lang en_US
langsupport --default en_US
keyboard us
mouse genericps/2
text
install
reboot
timezone --utc America/Los_Angeles
part
```


Kickstart File

◆ Packages section

```
%packages --ignoredeps --ignoremissing
@Base
PyXML
atlas
autofs
bc
chkrootkit
contrib-pexpect
contrib-pvfs-config
contrib-python-openssl
```



Kickstart File

◆ Post section

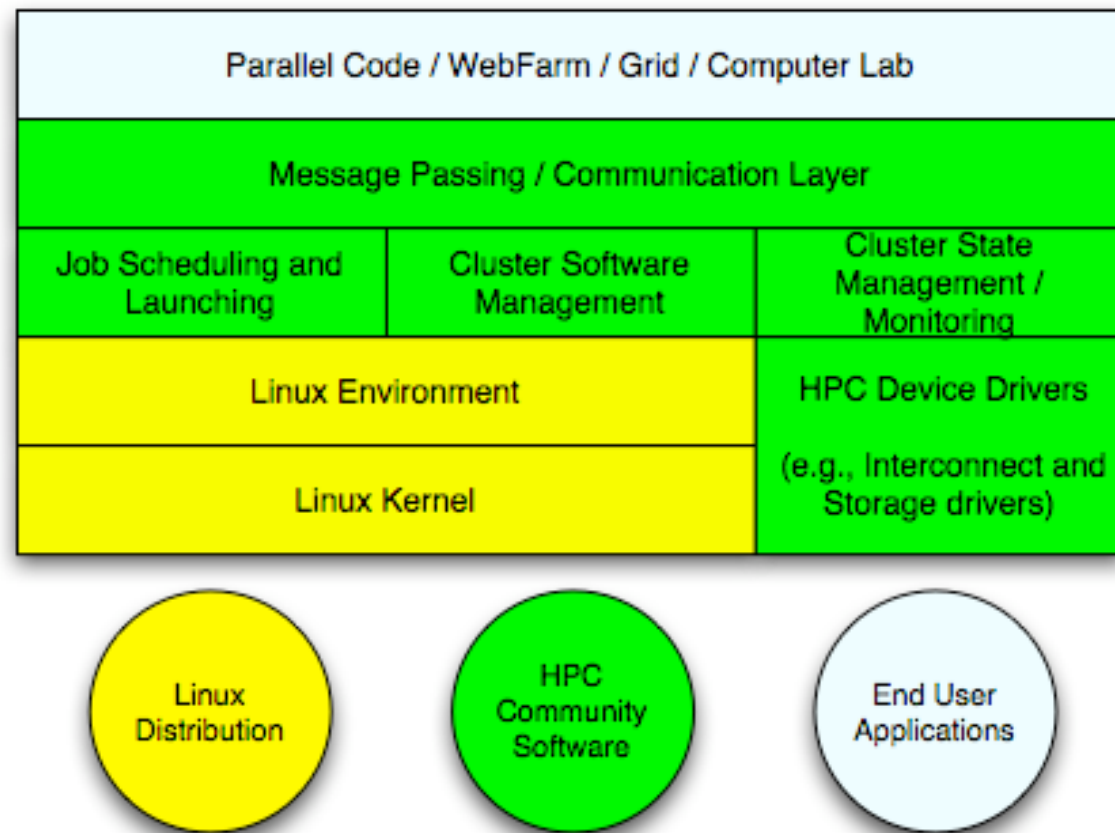
```
%post
```

```
cat > /etc/motd << 'EOF'  
Rocks Compute Node  
EOF
```

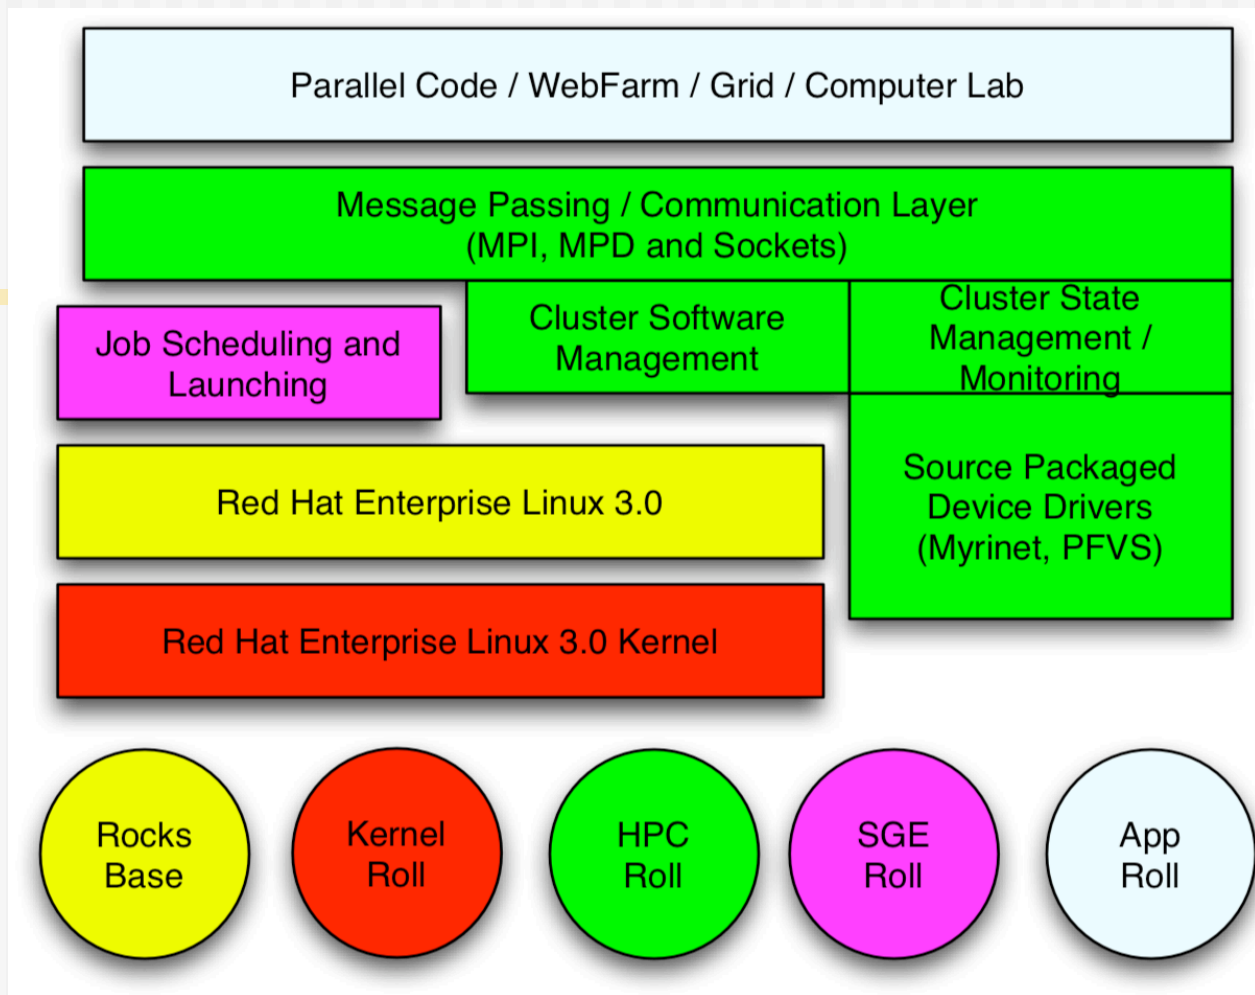


Rolls High-Level Description

Monolithic Software Stack



Rolls



◆ Dissecting the monolithic software stack

Rolls

PICK PACKAGES

- > COMBO #1: PREMIUM
- > COMBO #2: SPORT
- > COMBO #3: COLD WEATHER
- > NEXT STEP



Sport Package (\$1350)

CLICK IMAGE TO ADD THE SPORT PACKAGE TO YOUR LIST.

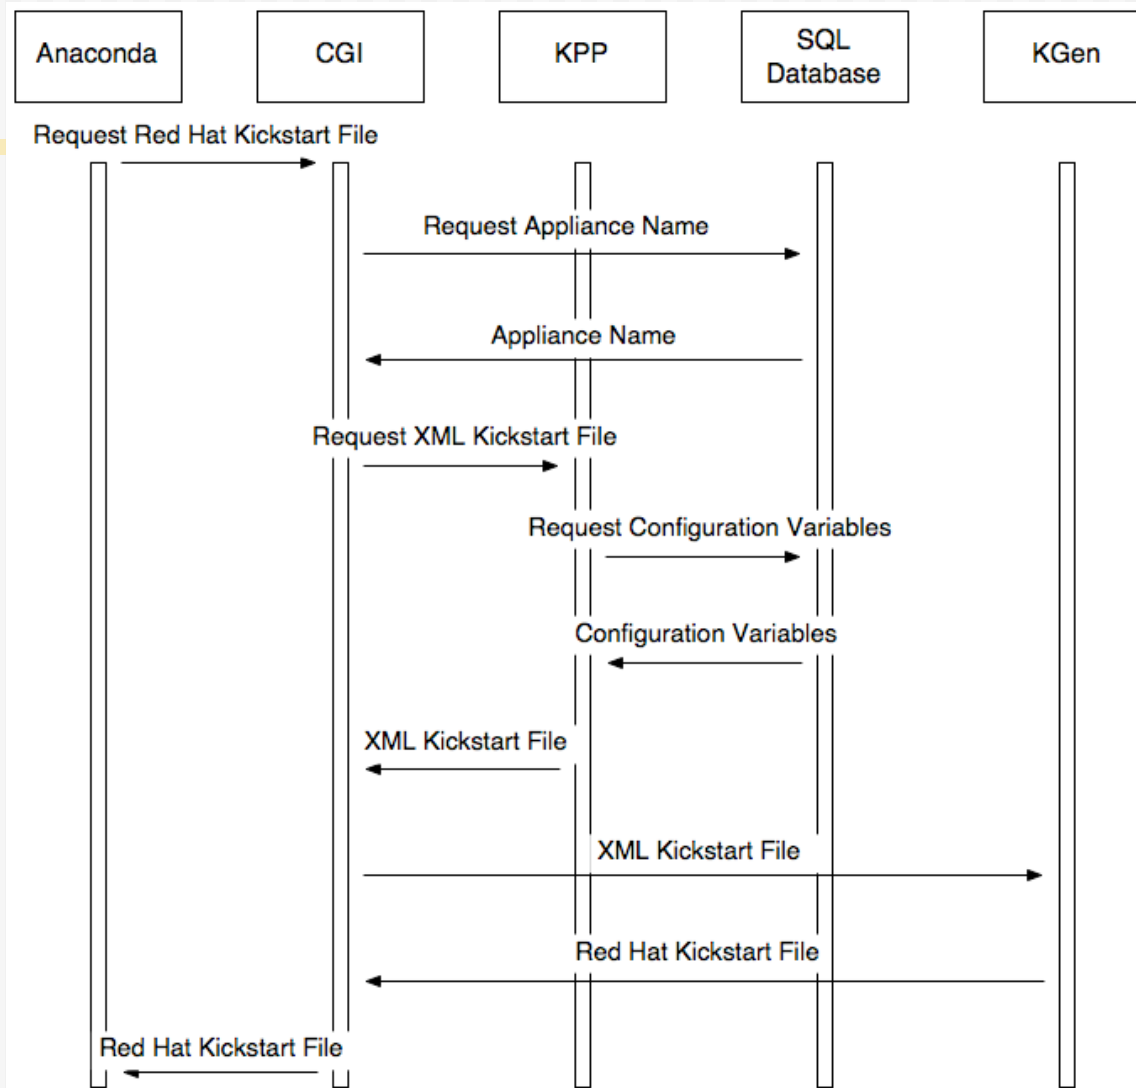
THE SPORT PACKAGE WILL ADD:

Dynamic stability control (DSC), bonnet stripes, xenon headlamps with powerwashers, front fog lamps, 17-inch alloy S-lite wheels with 205/45 R17 performance or all-season run-flat tires.

◆ Think of a roll as a “package” on a car




Getting A Kickstart File



Use Graph Structure to Dissect Distribution

- ◆ Use 'nodes' and 'edges' to build a customized kickstart file
- ◆ Nodes contain portion of kickstart file
 - ⇒ Can have a 'main', 'package' and 'post' section in node file
- ◆ Edges used to coalesce node files into one kickstart file

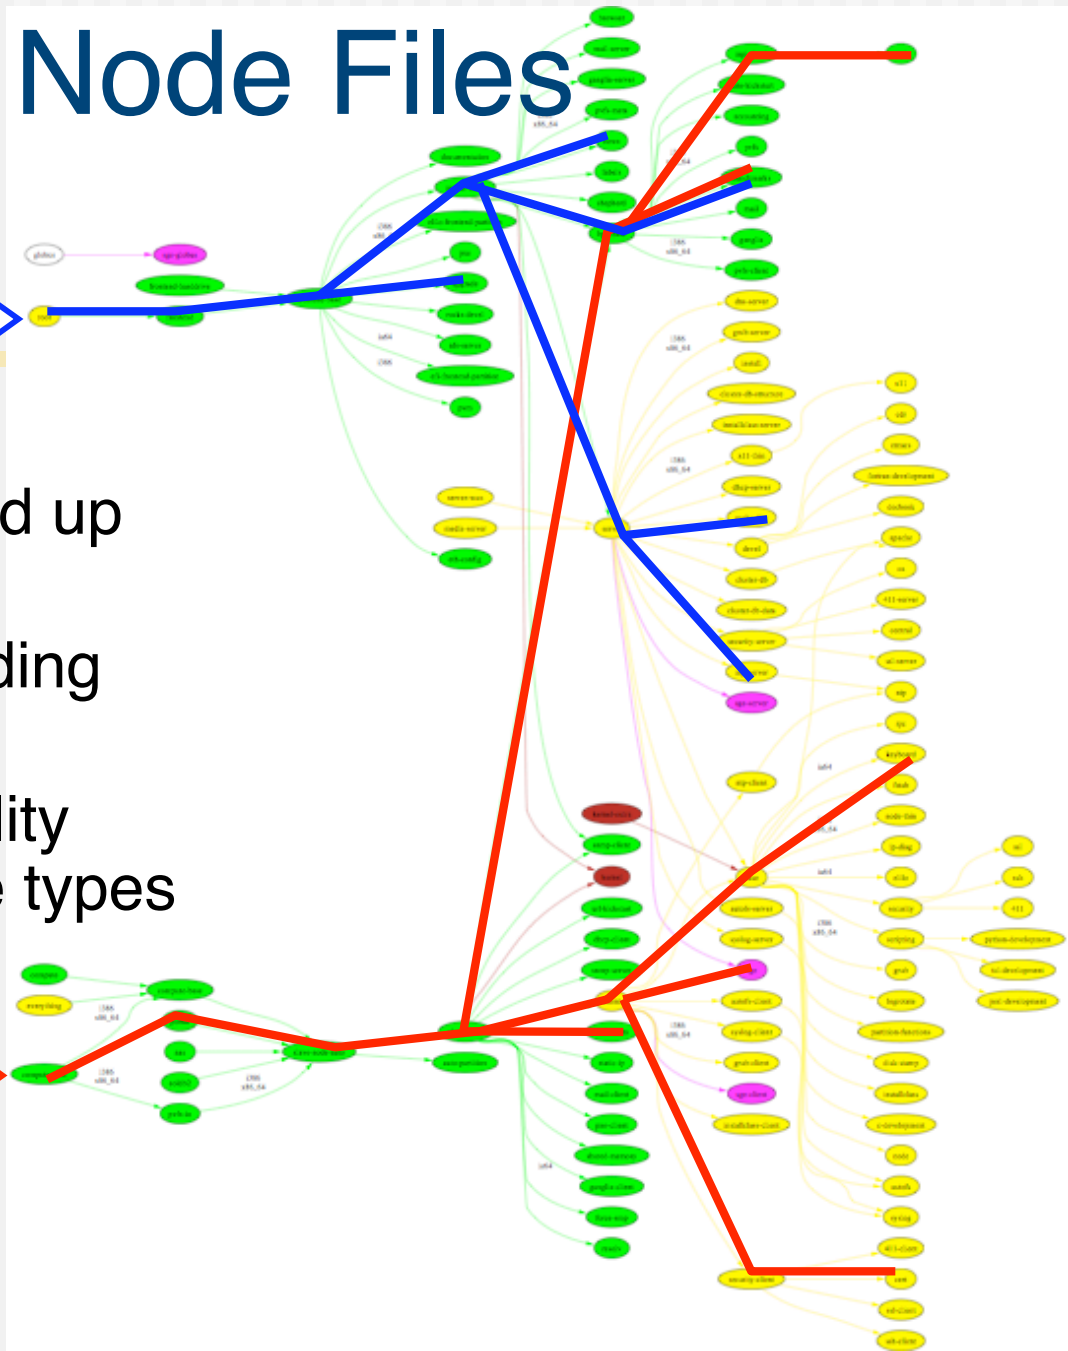


Coalescing Node Files

Frontend
Root



- ◆ Traverse a graph to build up a kickstart file
- ◆ Makes kickstart file building flexible
- ◆ Easy to share functionality between disparate node types

Compute
Root



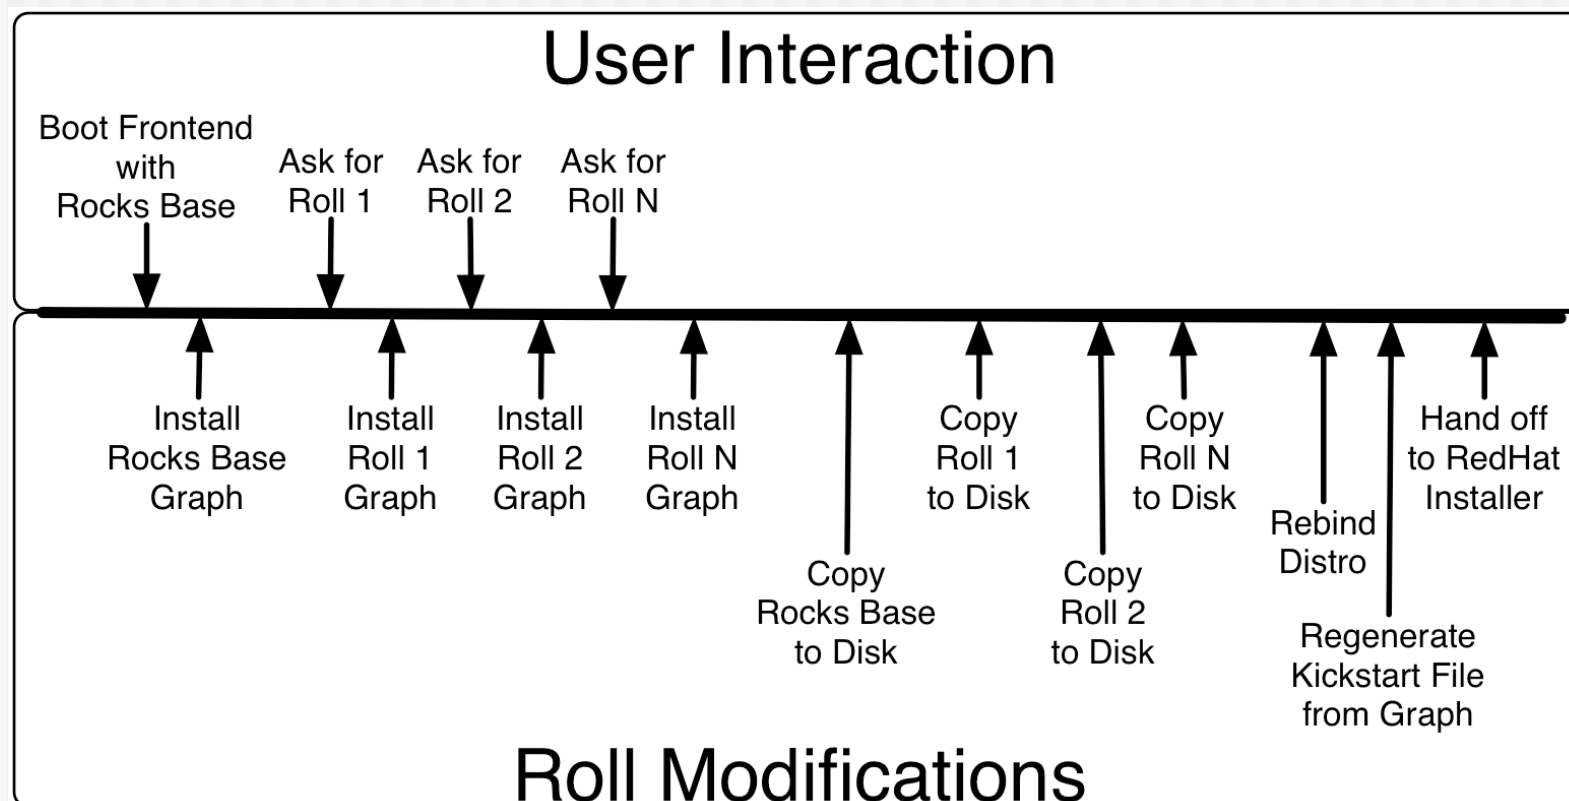
Why We Use A Graph

- ◆ A graph makes it easy to 'splice' in new nodes
- ◆ Each Roll contains its own nodes and splices them into the system graph file



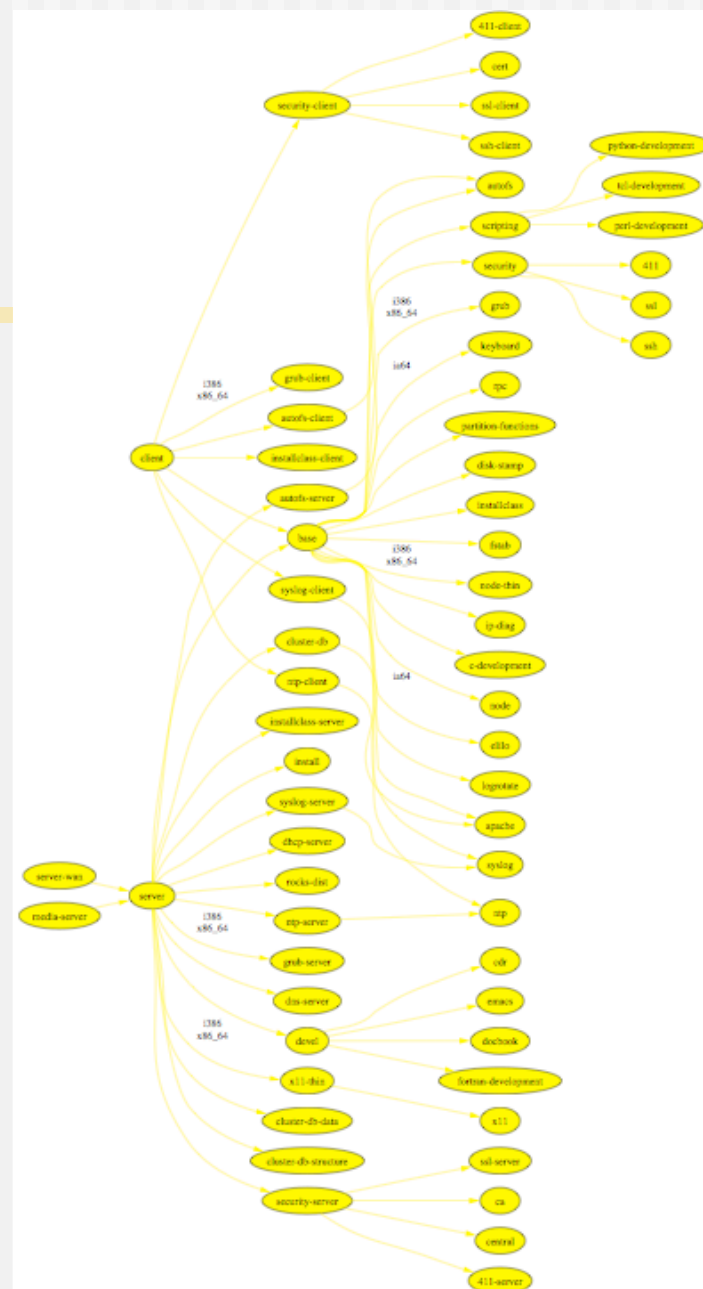


Rocks Extensions Installation Timeline



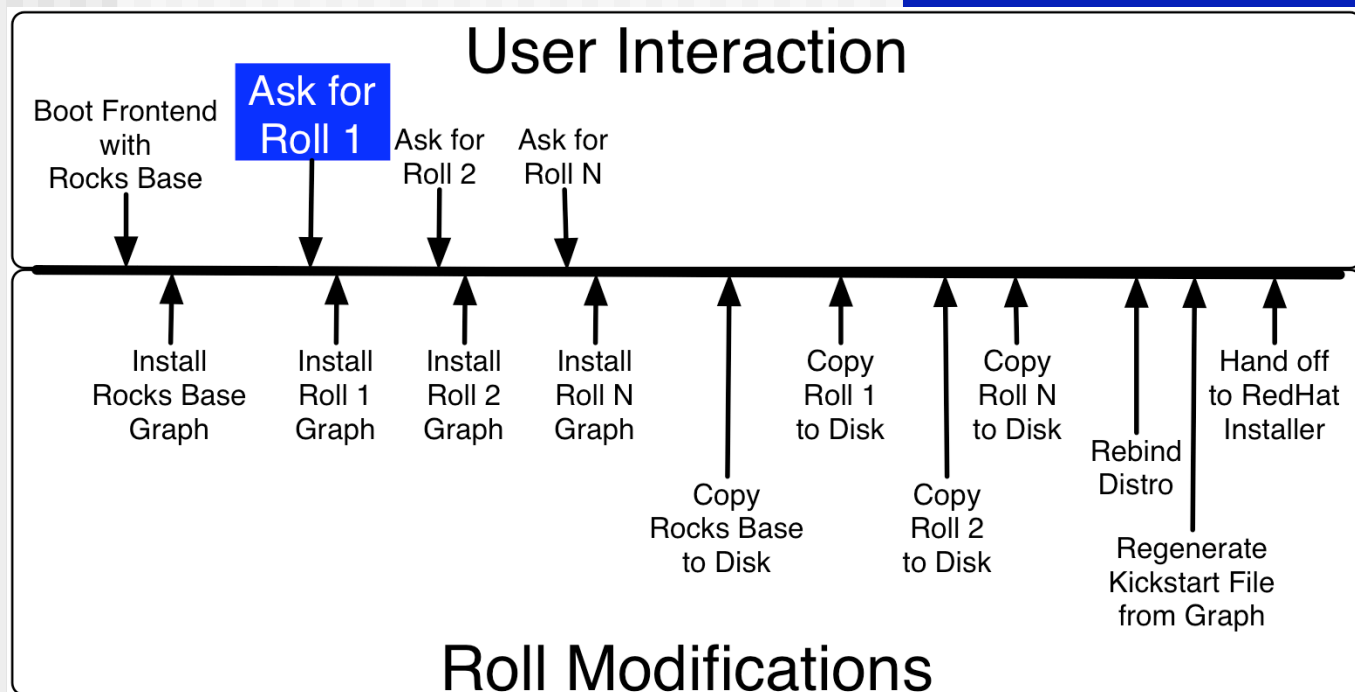



Install Rocks Base Graph






Anaconda Modified to Accept Rolls







Base + All Rolls



Anaconda Modified to Display New User Input Screens

Rocks v3.3.0 (Makalu) -- www.rocksclusters.org

Cluster Information

Fill in at least the FQDN

Fully Qualified Hostname:	Country:
<u>cluster.hpc.org</u>	<u>US</u>
Cluster Name: <u>mk</u>	Contact:
<u>Cluster</u>	<u>admin@cluster.hpc.org</u>
Organization:	URL:
<u>Hpc</u>	<u>http://cluster.hpc.org/</u>
Locality: <u>san.mk</u>	LatLong:
<u>San Diego</u>	<u>N32.87 W117.22</u>
State: <u>er: Updating graphs</u>	
<u>California</u>	

OK Back

<Tab>/<Alt-Tab> between elements | <Space> selects | <F12> next screen

Anaconda Modified to Display New User Input Screens

◆ How we do it:

- Place a shim in Anaconda to call our screens instead of the 'betanag' RedHat screen

```
index = 0
for key in installSteps:
    if key[0] == "betanag":
        break
    index = index + 1

installSteps[index] = ("rockswindows", ("id.rock", ))

# set list of user-defined windows
dispatch.skipStep("rockswindows", skip = 0)
stepToClasses["rockswindows"] = ("ksclass",
    tuple(rockswindows))
```

Anaconda Modified to Display New User Input Screens

◆ How you use it:

➔ In your XML file:

```
<installclass>
```

```
class CACLInfoWindow:
```

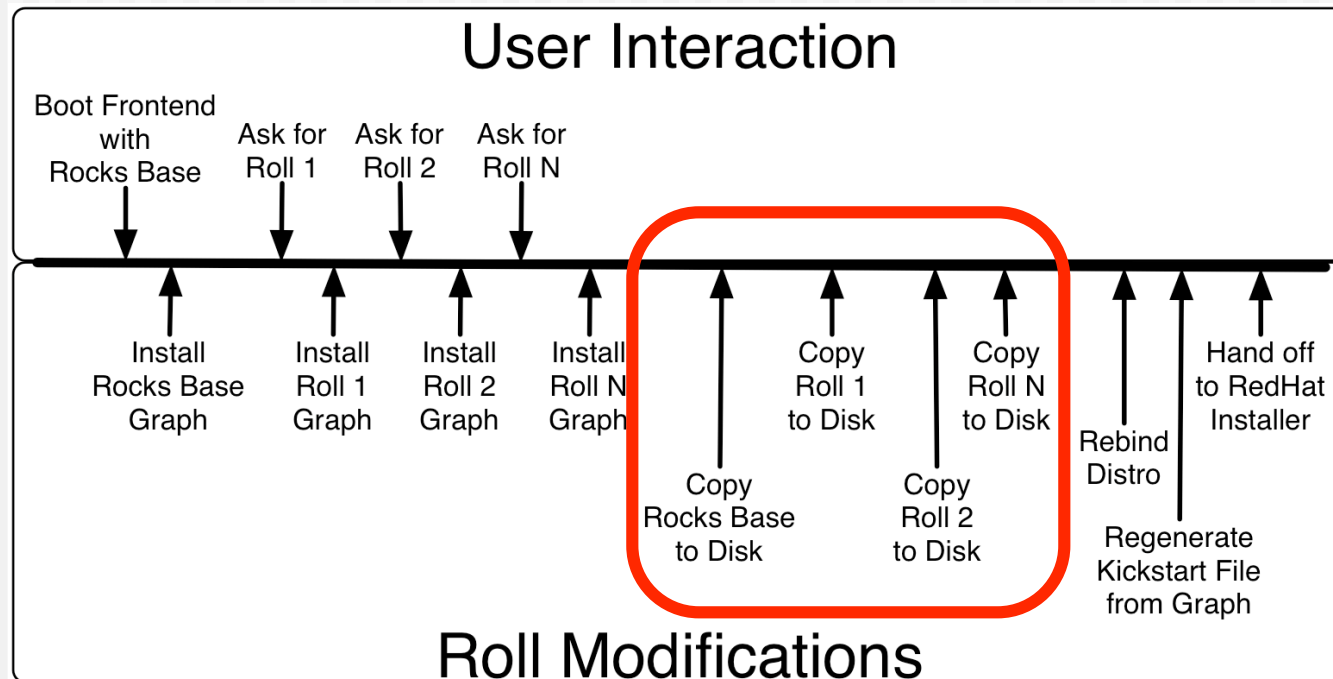
```
def __call__(self, screen, Info):
```

```
    bb = ButtonBar (screen, (TEXT_OK_BUTTON, ("Back","back")))
    toplevel = GridFormHelp (screen,
                             _("CACL Setup Information"), "CACLInfo", 1, 3)
    leftGrid = Grid(1,12)
    rightGrid = Grid(1,12)
    infoGrid = Grid(2,1)
    .
    .
```

```
addScreen("CACLInfoWindow")
```

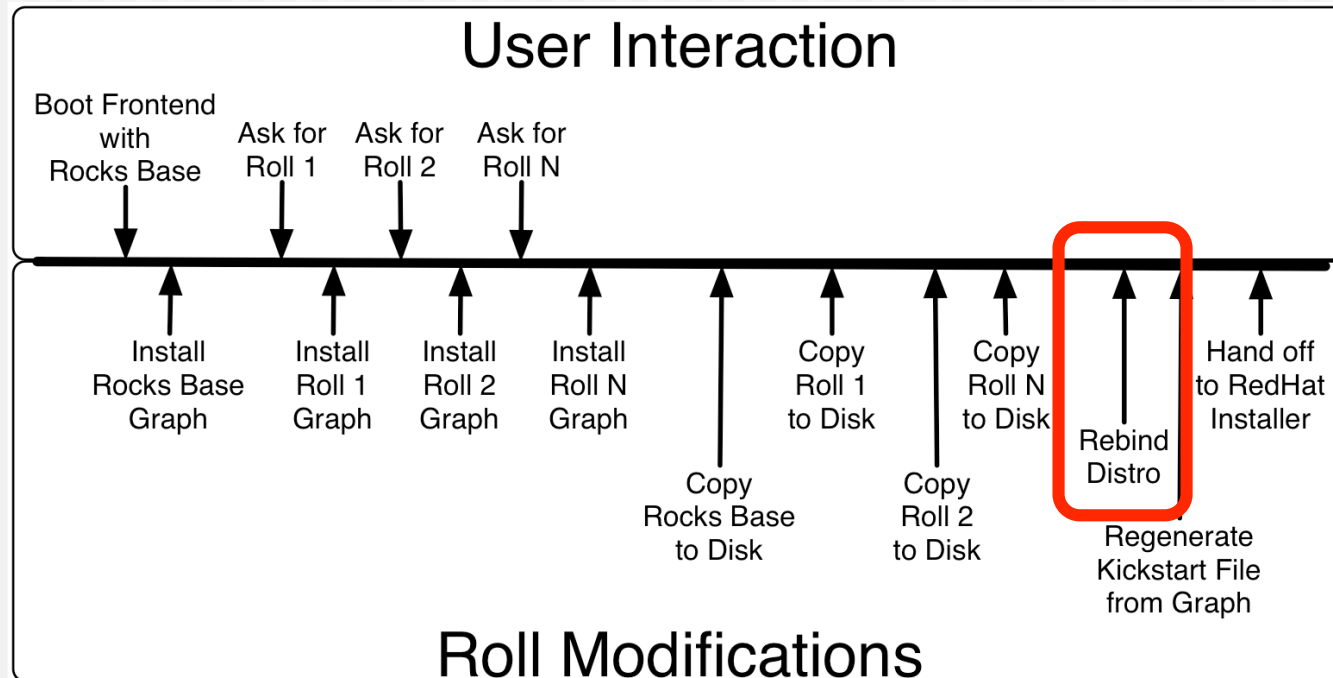
```
</installclass>
```

Copy Media To Local Disk



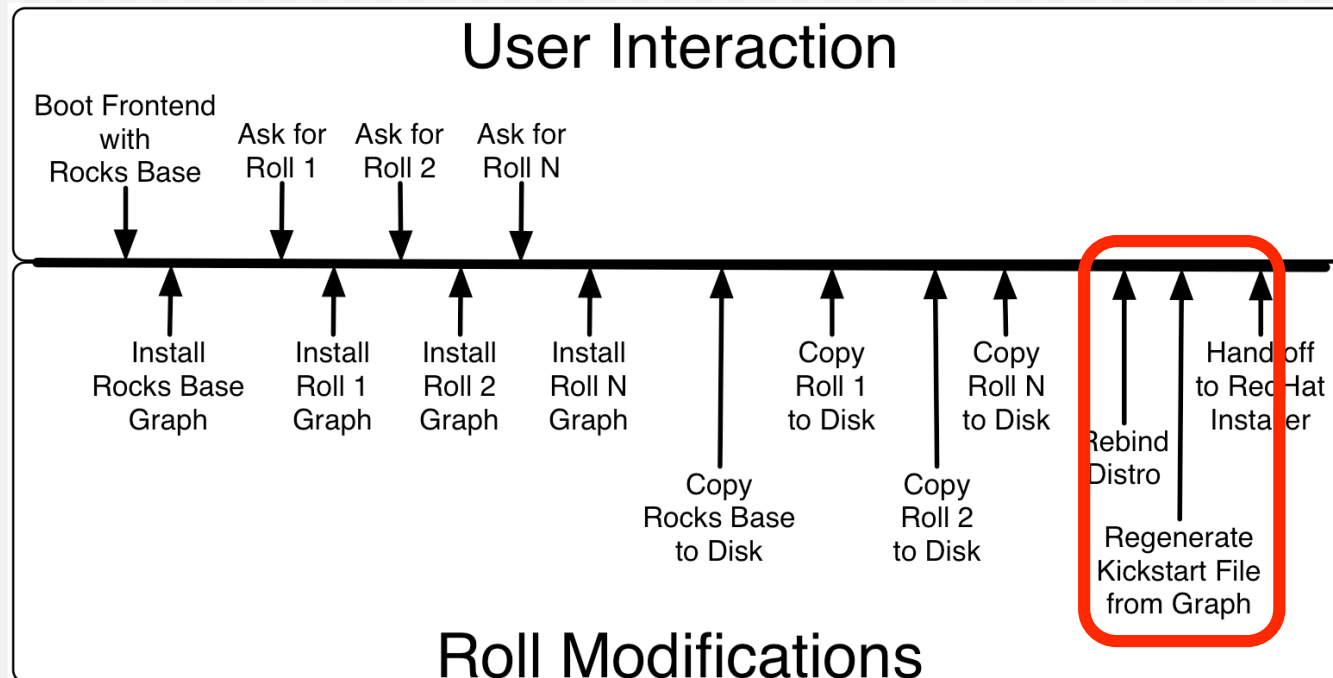
- ◆ Base and all user-supplied Rolls are copied to local disk
 - ➔ These packages are used to install compute nodes

Rebind Distro



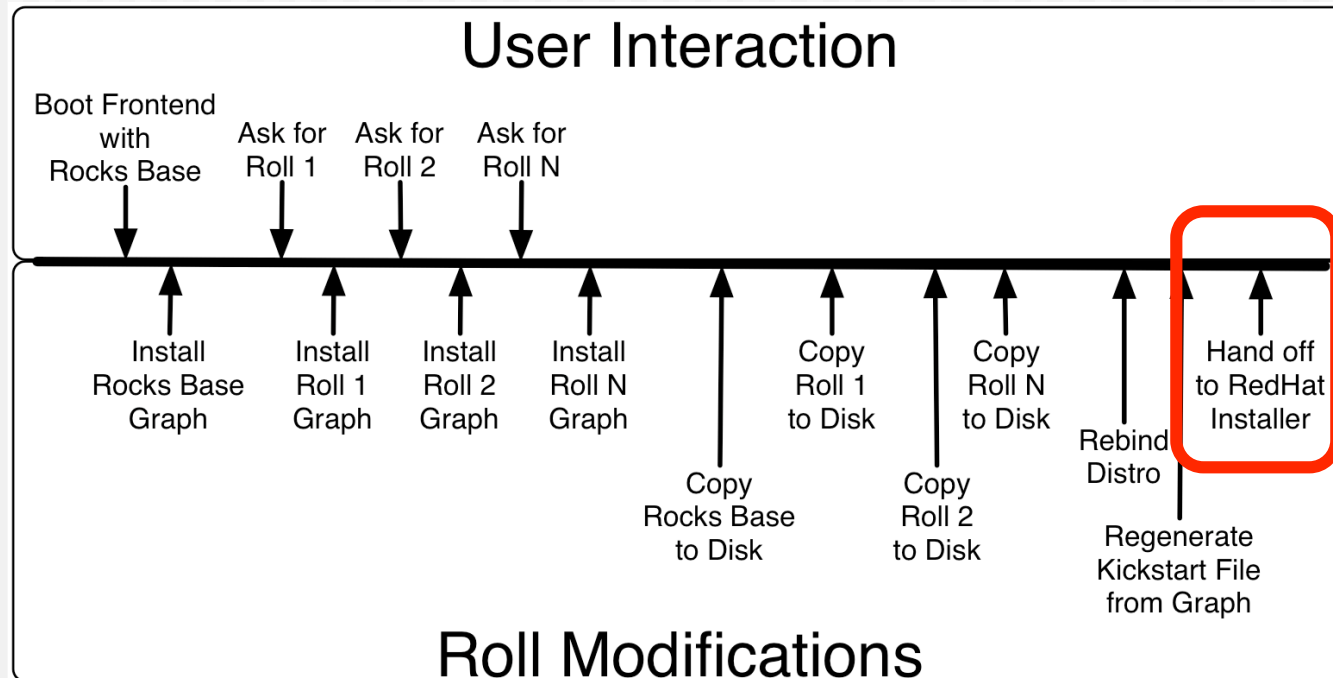
- ◆ Merge base with rolls into one RedHat-compliant distribution
 - ➡ This takes the dissected distro and tightly binds it
 - Note: We actually install the frontend off the local hard disk (not the CD media)

Rebuild the Kickstart File



- ◆ Traverse the final graph using the node 'frontend' as the root
 - ➔ Allows us to customize a frontend configuration at install time

Hand Off To RedHat



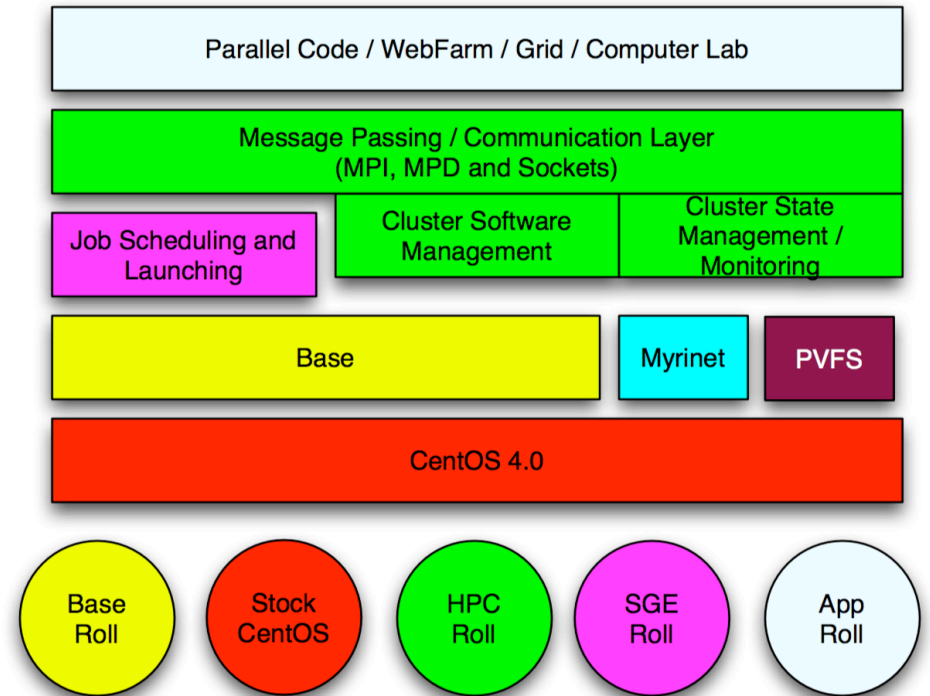
- ◆ Anaconda has no idea what hit it!
- ◆ The remainder of the installation looks like a standard RedHat installation (just with more packages and cluster-specific configuration)



Near Future


Rocks 4.0.0

- ◆ Everything is a Roll!



- ◆ Using CentOS 4.0 as base OS
 - ⇒ People have also successfully substituted Scientific Linux and RHEL 4 for CentOS
- ◆ When asked for roll, input stock CentOS CDs

Rocks 4.0.0



- ◆ Dissecting existing Rolls
 - ⇒ HPC roll is much smaller
 - E.g., Myrinet and PVFS are separate Rolls



Loader Modifications

Loader Modifications

- ◆ The first program that runs during a RedHat install is a C program called “loader”
- ◆ Performs low-level setup
 - ⇒ Loads drivers
 - ⇒ Configures network
 - ⇒ Downloads anaconda
 - ⇒ Gets kickstart file

Loader Modifications

- ◆ Make HTTP the default install method
 - ⇒ RedHat uses NFS as default

- ◆ Rationale
 - ⇒ Installation is read-only, don't need a file system
 - ⇒ HTTP traffic can be easily load balanced
 - ⇒ In terms of active development, more people are working on HTTP than NFS

Loader Modifications

◆ Robust kickstart file acquisition

- ⇒ 10 retries to get kickstart file
 - RedHat has only 1
- ⇒ NACK to throttle kickstart file acquisition
 - When load on frontend is high, the compute node is told to wait before next retry

◆ Rationale

- ⇒ The kickstart file is everything -- without it, a node is just a \$2,000 paperweight
- ⇒ NACK feature is for supporting large cluster reinstallations

Loader Modifications

◆ Watchdog

- ⇒ If can't get kickstart file or if there is an error during the installation, reboot
 - This will restart the installation
 - RedHat just halts

◆ Rationale

- ⇒ Again, the kickstart file is everything

Loader Modifications

- ◆ Network-based frontend installations
 - ⇒ In Rocks lingo: a “central” install

- ◆ Rationale
 - ⇒ The “CD dance” during installation is not optimal
 - ⇒ Needed to support grids of clusters from a central place
 - ⇒ Huge benefit for development
 - Don’t have to burn CDs just to test code changes

Loader Modifications

- ◆ Secure kickstart
 - ➔ Added HTTPS support
- ◆ Rationale
 - ➔ Needed for support of network-based frontend installations (“central” installs)
 - Don’t want the root password for the frontend sent over the network in the clear!
 - ➔ Useful for compute nodes that are installed over a public network

Loader Modifications

- ◆ Support adding compute node to any ethernet interface
 - ➔ The first interface that receives a kickstart file, is anointed 'eth0'
- ◆ Rationale
 - ➔ Email reduction
 - We got lots of email from people who plugged their ethernet cable into the “wrong” port
 - ➔ Even the Three Stooges can plug in the cables to the compute nodes!

Loader Modifications

◆ Bug Fixes

- ➔ Added support for multiple CD drives
- ➔ A couple stack overflow problems

◆ Rationale

- ➔ Without the fixes, the installer halts



In The Next Session

Graph details