



# Rocks 101

---

Rocks-A-Palooza I  
Track 1  
Session II



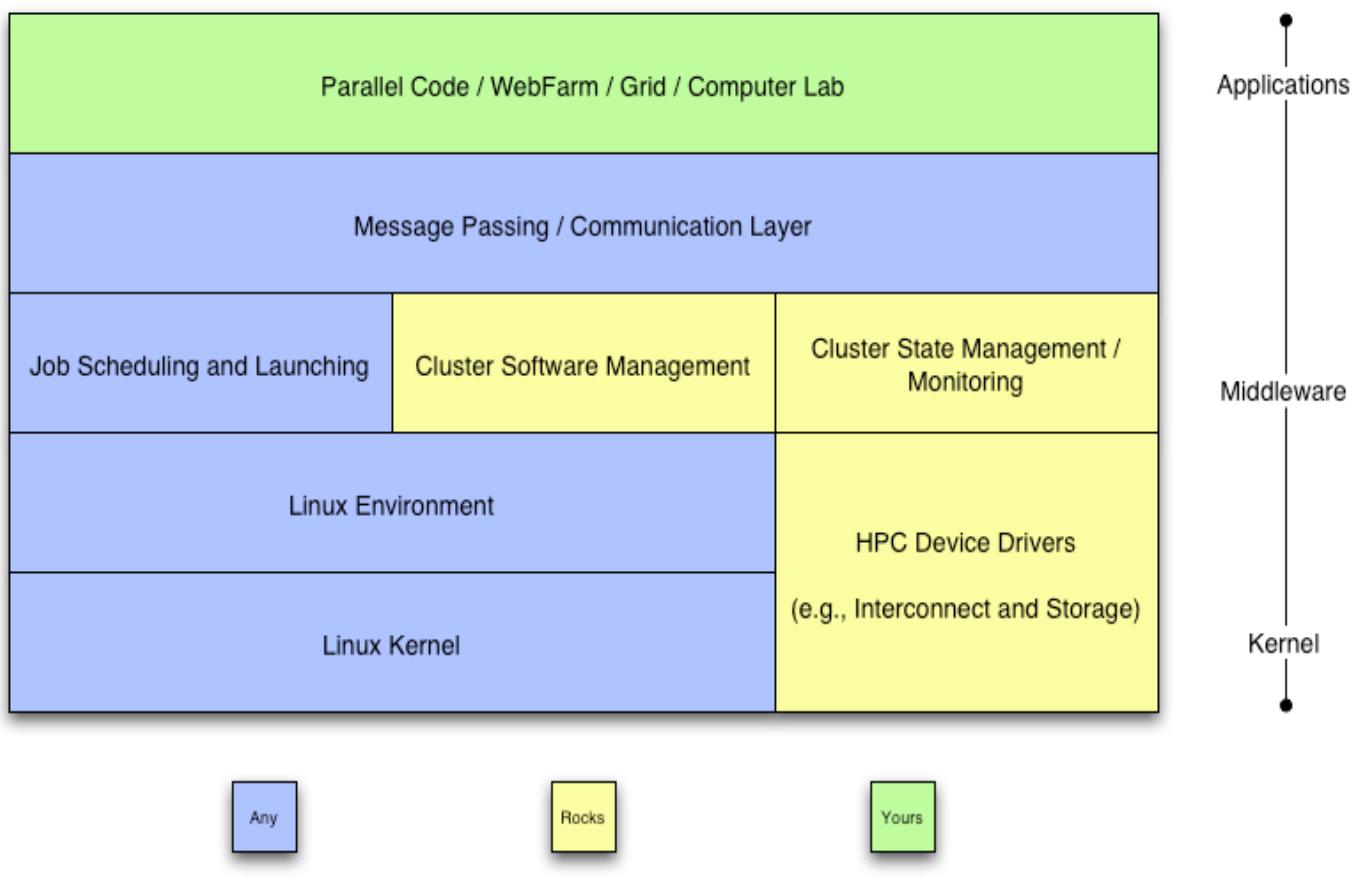
**ROCKS**

# Overview of Rocks

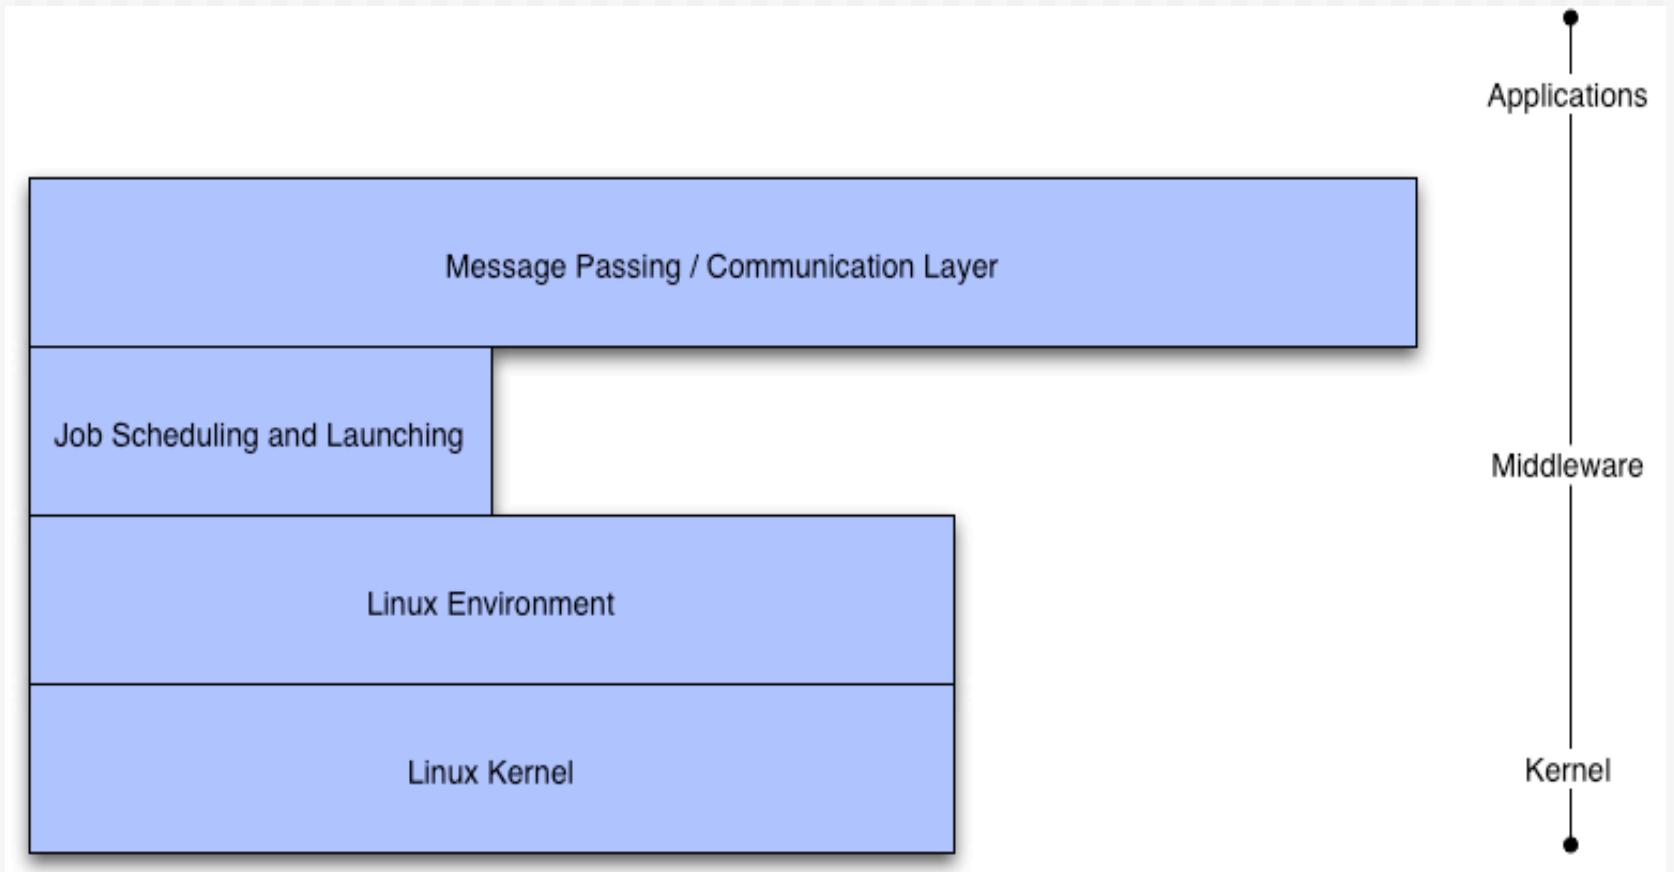
---

The Rocks software stack

# Cluster Software Stack



# Common to Any Cluster





# Red Hat

- ◆ Enterprise Linux 3.0
  - ➲ Recompiled from public SRPMS, including errata updates (source code)
  - ➲ No license fee required, redistribution is also fine
  - ➲ Recompiled for all CPU types (x86, Opteron, Itanium)
  - ➲ *Rocks 4.0 is based on RHEL 4.0 (Centos, or RHEL)*
- ◆ Standard Red Hat Linux kernel
  - ➲ No Rocks added kernel patches
- ◆ No support for other distributions
  - ➲ Red Hat is the market leader for Linux
    - In the US
    - And becoming so in Europe
  - ➲ Trivial to support any Anaconda-based system
  - ➲ Others would be harder, and require vendor support (SuSe ~ 12 months work)
- ◆ Excellent support for automated installation
  - ➲ Scriptable installation (Kickstart)
  - ➲ Very good hardware detection

# Dell Invests in Red Hat

## Michael Dell puts \$99.5M in Red Hat

Billionaire chairman of No. 1 PC maker places big bet on Microsoft competitor.

May 10, 2005: 1:41 PM EDT

**NEW YORK (CNN/Money)** - Red Hat is getting a \$99.5 million boost from Michael S. Dell, billionaire founder and chairman of Dell Inc., according a regulatory filing.

Through his private investment firm, MSD, Dell bought the largest share of \$600 million in debentures offered by the software developer in January 2004, a Securities Exchange Commission filing showed.

Red Hat's main product, the Linux operating system for PCs, is a direct competitor to Microsoft's Windows. The Raleigh, N.C.-based company also provides support services for "open source" technology, which is software developed by communities of programmers for free use.

[Dell \(Research\)](#) is the nation's largest PC maker.

Debentures are similar to bonds in that the issuer promises a fixed return for a stated period of time on the investment.

In the case of a public company, a debenture can also be converted into shares or equity. ■



COURTESY: DELL COMPUTER

Michael Dell, billionaire chairman of Dell Inc., has given Red Hat a \$99.5M injection.



## Tech

[SEE MORE STORIES](#)

advertiser links

[what's this?](#)

### [Accounting Research Manager](#)

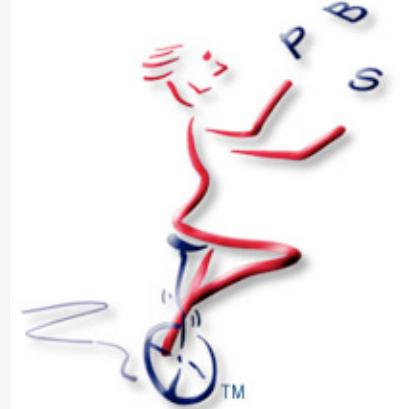
Find insightful interpretations on GAAP and Securities and Exchange Commission...

[www.accountingresearchmanager.com](http://www.accountingresearchmanager.com)

### [Securities Exchange Commission](#)

# Batch Systems

- ◆ Portable Batch System and Maui
  - ➲ Long time standard for HPC queuing systems
  - ➲ Maui provides backfilling for high throughput
  - ➲ PBS/Maui system can be fragile and unstable
  - ➲ Multiple code bases:
    - PBS
    - OpenPBS
    - PBSPro
    - Scalable PBS
- ◆ Sun Grid Engine
  - ➲ Rapidly becoming the new standard
  - ➲ Integrated into Rocks by Scalable System
    - See Najib and Sivaram
  - ➲ Now the default scheduler for Rocks
  - ➲ Robust and dynamic
  - ➲ Currently 5.3, moving to 6.0 when out of Beta



# Communication Layer

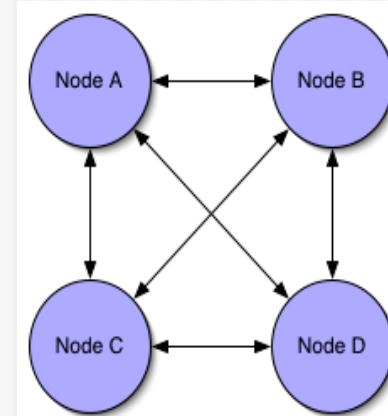
---

- ◆ None
  - ➲ “Embarrassingly Parallel”
- ◆ Sockets
  - ➲ Client-Server model
  - ➲ Point-to-point communication
- ◆ MPI - Message Passing Interface
  - ➲ Message Passing
  - ➲ Static model of participants
- ◆ PVM - Parallel Virtual Machines
  - ➲ Message Passing
  - ➲ For Heterogeneous architectures
  - ➲ Resource Control and Fault Tolerance

# Sockets are low level

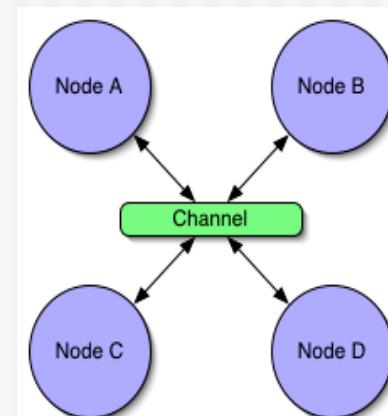
- ◆ Sockets

- ↳ Point-to-Point
- ↳  $N$  machines =  $(n^2 - n)/2$  connections
- ↳ 1, 3, 6, 10, 15, ...



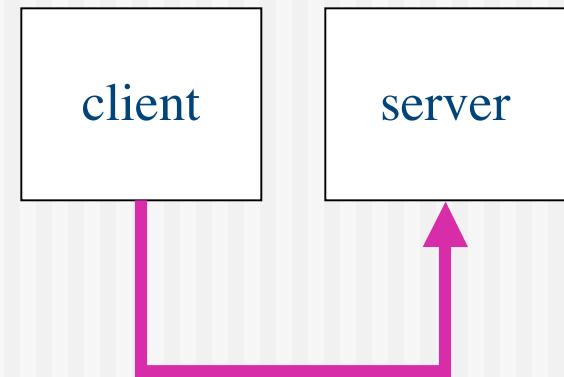
- ◆ MPI/PVM

- ↳ Shared virtual channel
- ↳ Implementation could be sockets
- ↳ Easier to program



# Sockets

- ◆ Open an endpoint
- ◆ Specify IP address and port
- ◆ Send / receive messages
  - ➲ If TCP, only point-to-point messages
  - ➲ If UDP, option of point-to-point or multicast (broadcast)
- ◆ Shutdown connection



# High-level TCP Example

```
/*
 * SERVER CODE
 */
fd = socket();
.
.
.
saddr.s_addr      = INADDR_ANY;
saddr.port        = 1234;
bind(fd, &saddr);
listen(fd);
accept(fd);
.
.
.
read(fd, buffer, size);
.
.
.
close(fd);

/*
 * CLIENT CODE
 */
fd = socket();
.
.
.
saddr.s_addr      = gethostbyname("c0-0");
saddr.port        = 1234;
.
.
.
write(fd, buffer, size);
.
.
.
close(fd);
```

# Challenges with Sockets

- ◆ TCP

- ↳ Reliable, but byte oriented
- ↳ Need to write code to send and receive *packets* (at the application level)

- ◆ UDP

- ↳ Unreliable
- ↳ Need to write code to reliably send packets

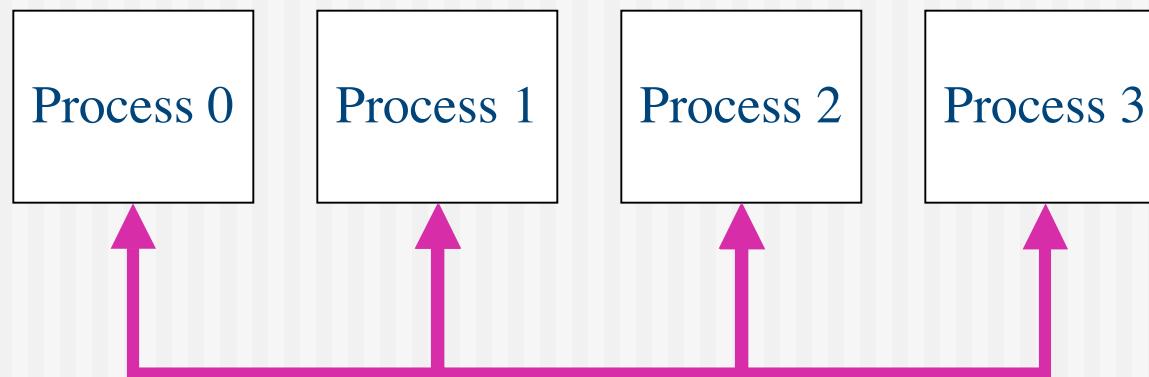
# MPI

---

- ◆ Message Passing Interface
- ◆ De facto standard for message passing
  - ⇒ Runs over many CPU architectures and many communication substrates
- ◆ There are (and were) lots of good messaging libraries
  - ⇒ But, MPI is the most pervasive
  - ⇒ Developed a practical, portable, efficient and flexible standard
  - ⇒ In development since 1992

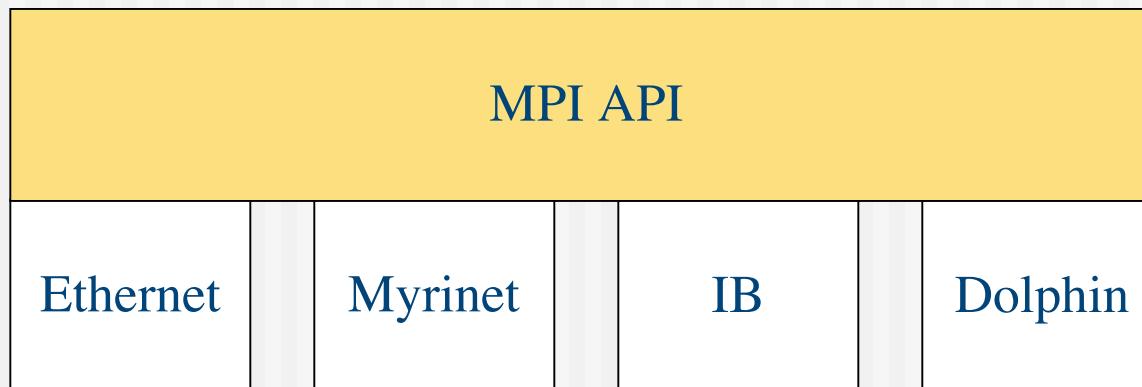
# MPI

- ◆ Explicitly move data like sockets, but virtualizes the endpoints
  - ⇒ Remote endpoints addressed by integer 0, 1, ..., n
- ◆ Primitives to support point-to-point and broadcast



# MPI

- ◆ Single interface to pass messages over many communication substrates



# High-level MPI Example

```
MPI_Init();
.
.
.
MPI_Comm_rank(&my_mpi_id);
.
.
.
Remote_mpi_id = 1
MPI_Send(send_buffer, buf_size, remote_mpi_id)
.
.
.
MPI_Recv(recv_buffer, buf_size, remote_mpi_id)
.
.
.
MPI_Finalize()
```

# Challenges with MPI

---

- ◆ If a node fails, no easy way to reconfigure and route around the problem
  - ⇒ Basically, your program stops

# Compile

## ◆ MPICH with GNU Compilers and Ethernet

### **Compiler Path**

C: /opt/mpich/ethernet/gcc/bin/mpicc  
C++: /opt/mpich/ethernet/gcc/bin/mpicxx  
F77: /opt/mpich/ethernet/gcc/bin/mpif77

## ◆ MPICH with GNU Compilers and Myrinet

### **Compiler Path**

C: /opt/mpich/myrinet/gcc/bin/mpicc  
C++: /opt/mpich/myrinet/gcc/bin/mpicxx  
F77: /opt/mpich/myrinet/g77/bin/mpif77

# Compile



## ◆ MPICH with Intel Compilers and Ethernet

### **Compiler Path**

C:	/opt/mpich/ethernet/ecc/mpicc
C++:	/opt/mpich/ethernet/ecc/mpiCC
F77:	/opt/mpich/ethernet/ecc/mpif77
F90:	/opt/mpich/ethernet/ecc/mpif90

## ◆ MPICH with Intel Compilers and Myrinet

### **Compiler Path**

C:	/opt/mpich/myrinet/ecc/mpicc
C++:	/opt/mpich/myrinet/ecc/mpiCC
F77:	/opt/mpich/myrinet/efc/mpif77
F90:	/opt/mpich/myrinet/efc/mpif90

# PVM

---

- ◆ Parallel Virtual Machines v3.4.3
  - ➲ Message passing interface for heterogeneous architectures
    - Supports over 60 variants of UNIX
    - Supports Windows NT
  - ➲ Resource control and meta computing
  - ➲ Fault tolerance
  - ➲ <http://www.csm.ornl.gov/pvm/>

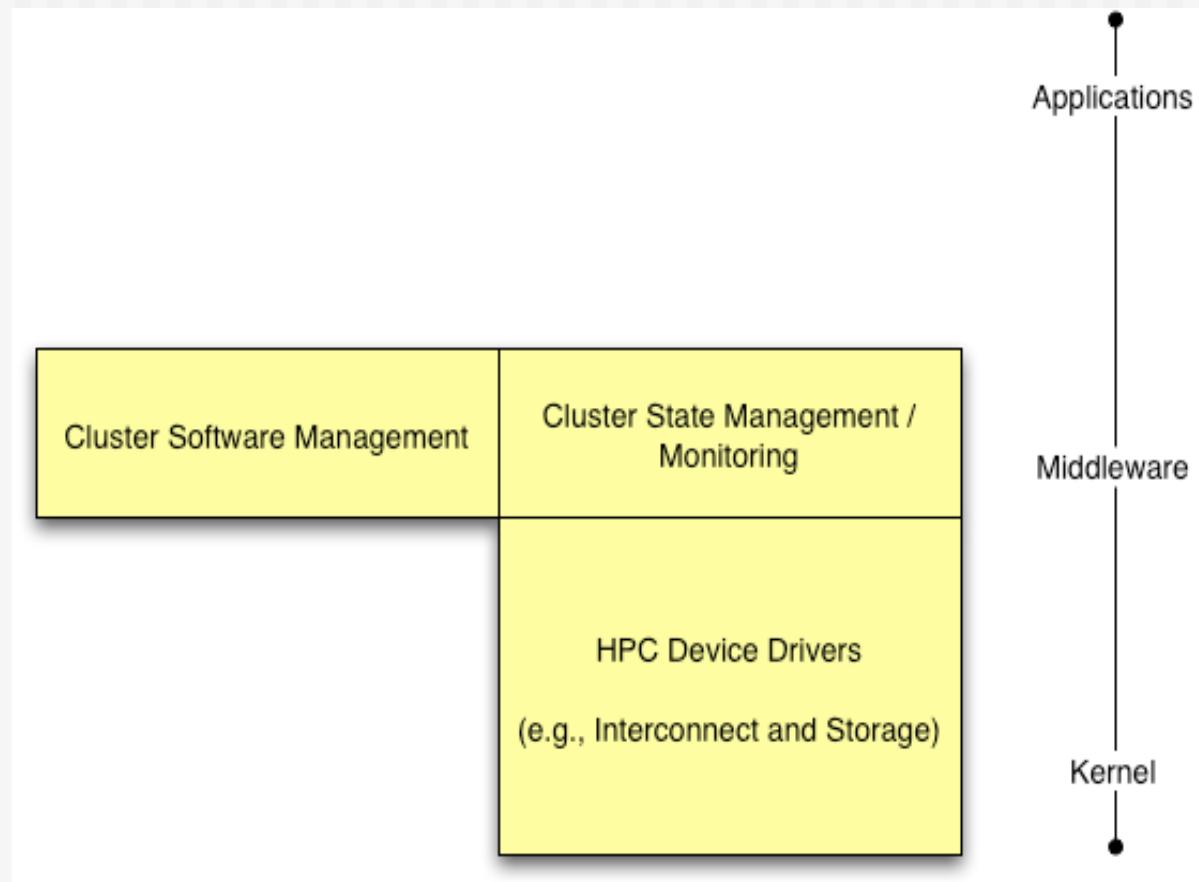
# NFS

- ◆ User account are served over NFS
  - ➲ Works for small clusters (<= 128 nodes)
  - ➲ Will not work for large clusters (>1024 nodes)
  - ➲ NAS is better than Linux
    - Rocks uses the Frontend machine to server NFS
    - We have deployed NAS on several clusters
- ◆ Applications are not served over NFS
  - ➲ /usr/local/ does not exist
  - ➲ All software is installed locally from RPM

# Open SSH

- ◆ Replaces Telnet, Rsh
  - ➲ Cryptographically strong authentication and encryption
  - ➲ Forwards X11 connections (no more \$DISPLAY)
- ◆ Rocks uses SSH
  - ➲ Mpirun
  - ➲ Cluster-fork
- ◆ Ssh-agent
  - ➲ Manager for SSH keys
  - ➲ ssh-agent \$SHELL

# Rocks Cluster Software



# SNMP

---

- ◆ Enabled on all compute nodes
- ◆ Great for point-to-point use
  - ⇒ Good for high detail on a single end-point
  - ⇒ Does not scale to full cluster wide use
- ◆ Supports Linux MIB
  - ⇒ Uptime, Load, Network statistics
  - ⇒ Install Software
  - ⇒ Running Processes

# Syslog

- ◆ Native UNIX system event logger
  - ➲ Logs events to local dist
    - /var/log/message
    - Rotates logs daily, eventually historic data is lost
  - ➲ Forwards all message to the frontend
- ◆ Scalable
  - ➲ Can add additional loghosts
  - ➲ Can throttle verbosity of loggers
- ◆ Uses
  - ➲ Predicting hardware and software failures
  - ➲ Post Mortem on crashed nodes
  - ➲ Debugging System startup

# eKV

---

- ◆ Remotely Interact with Installation
  - ⇒ Initial kickstart
  - ⇒ Re-Installation
- ◆ Shoot-node
  - ⇒ Reinstall OS and brings up eKV
- ◆ eKV
  - ⇒ Ssh to node while it is installing
  - ⇒ See the console output over Ethernet

# Cluster State Management

- ◆ Static Information

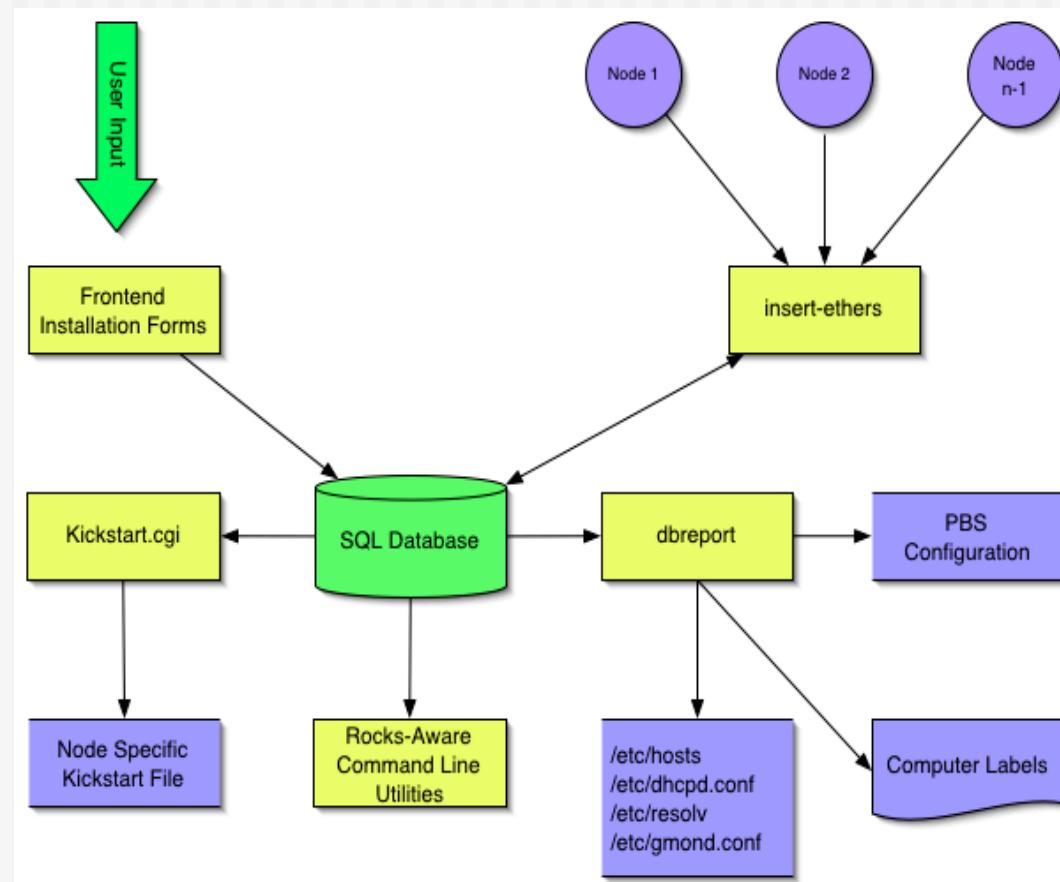
- ↳ Node addresses
- ↳ Node types
- ↳ Site-specific configuration

- ◆ Dynamic Information

- ↳ CPU utilization
- ↳ Disk utilization
- ↳ Which nodes are online

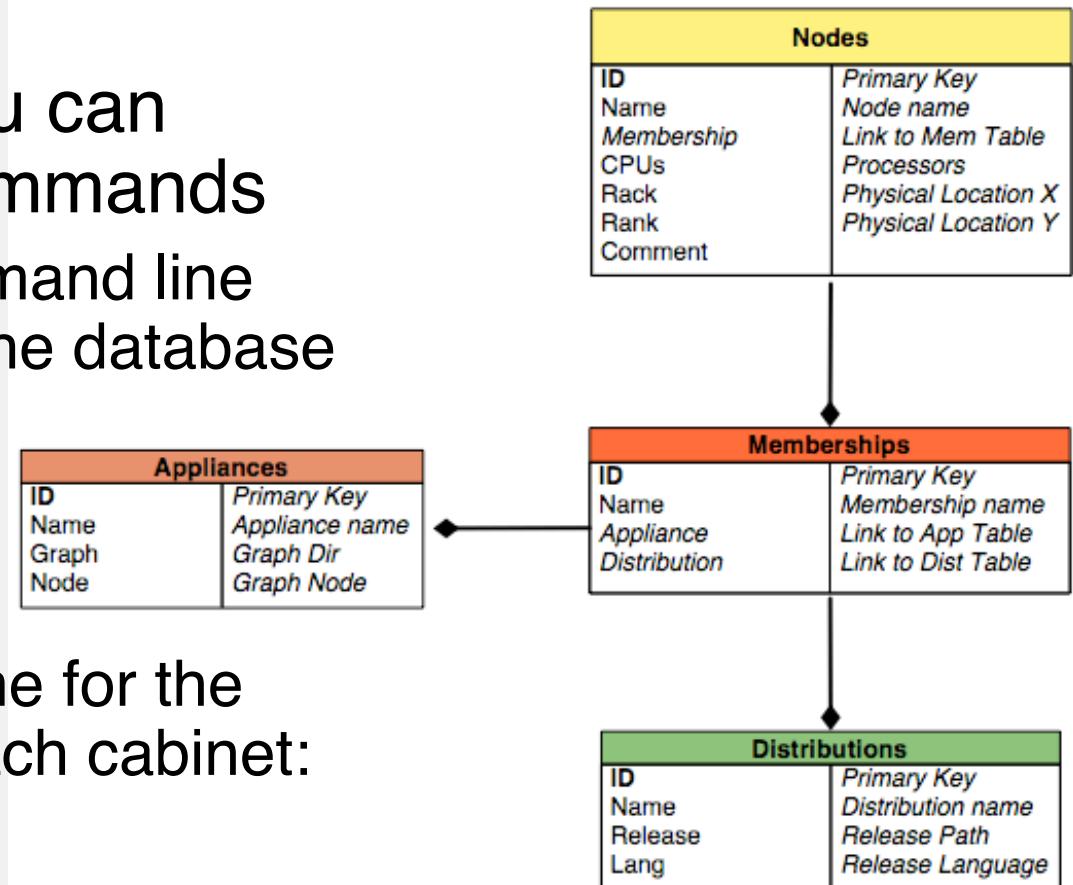


# Cluster Database



# Node Info Stored In A MySQL Database

- ◆ If you know SQL, you can execute powerful commands
  - ➲ Rocks-supplied command line utilities are tied into the database

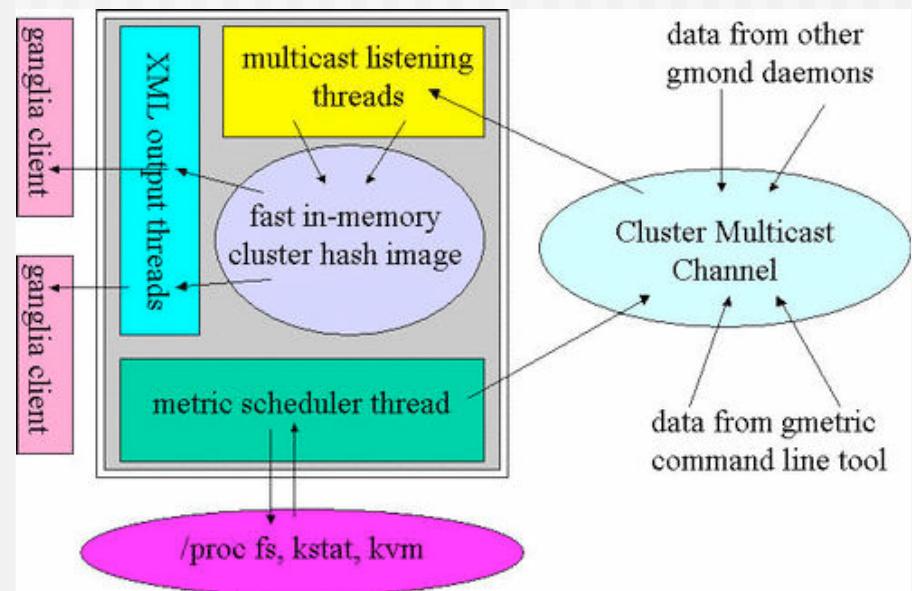


- ➲ E.g., get the hostname for the bottom 8 nodes of each cabinet:

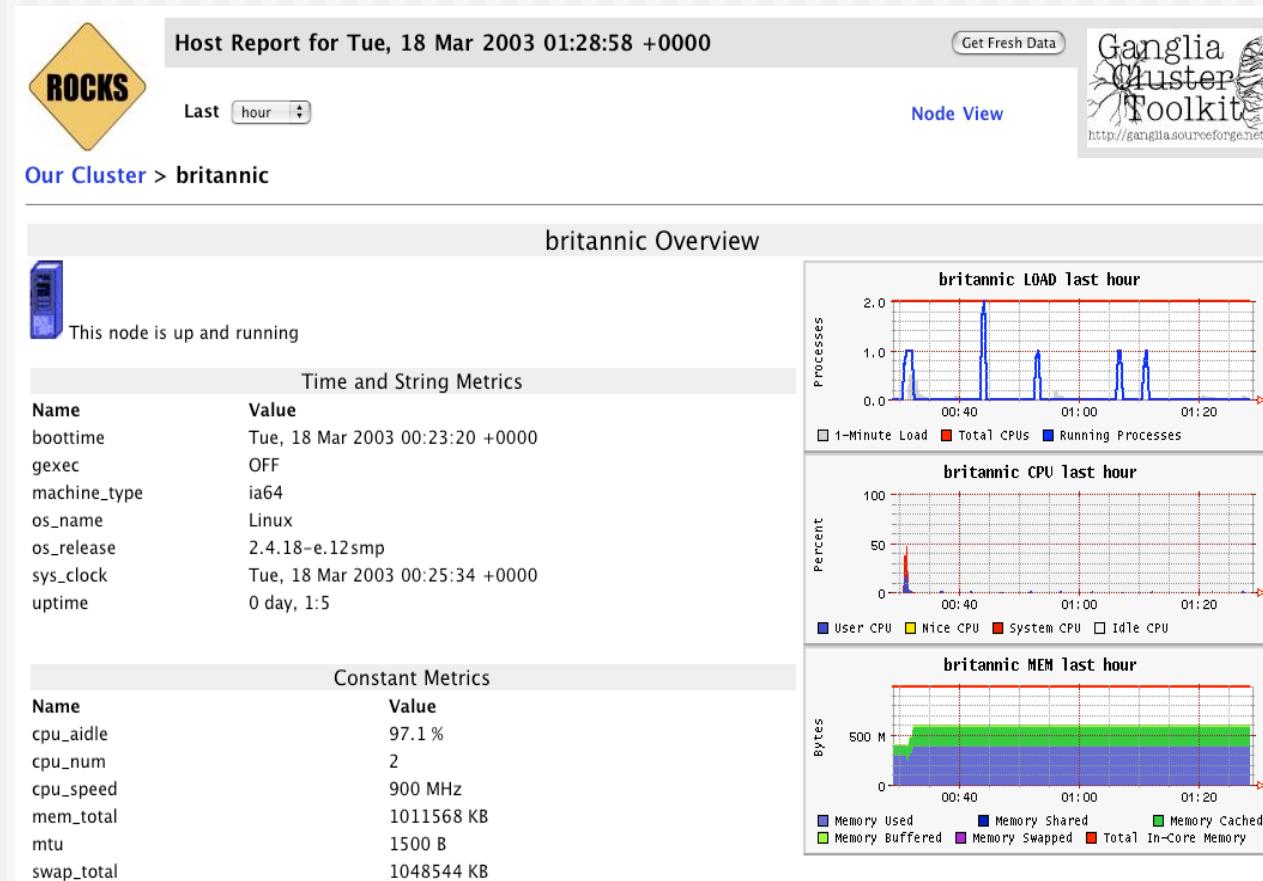
```
# cluster-fork --query="select name from nodes where rank<8" hostname
```

# Ganglia (or SCMSWeb / SCE Roll)

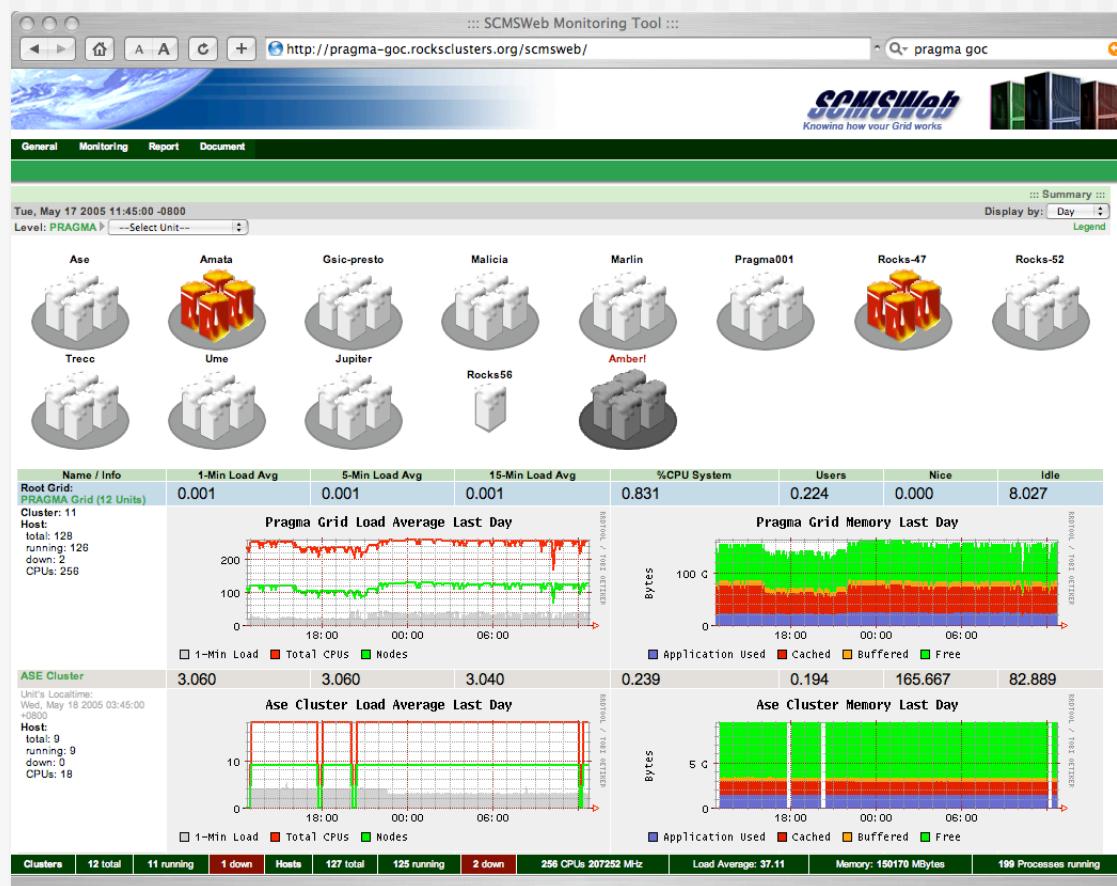
- ◆ Scalable cluster monitoring system
  - Based on ip multi-cast
  - Matt Massie, et al from UCB
  - <http://ganglia.sourceforge.net>
- ◆ Gmon daemon on every node
  - Multicasts system state
  - Listens to other daemons
  - All data is represented in XML
- ◆ Ganglia command line
  - Python code to parse XML to English
- ◆ Gmetric
  - Extends Ganglia
  - Command line to multicast single metrics



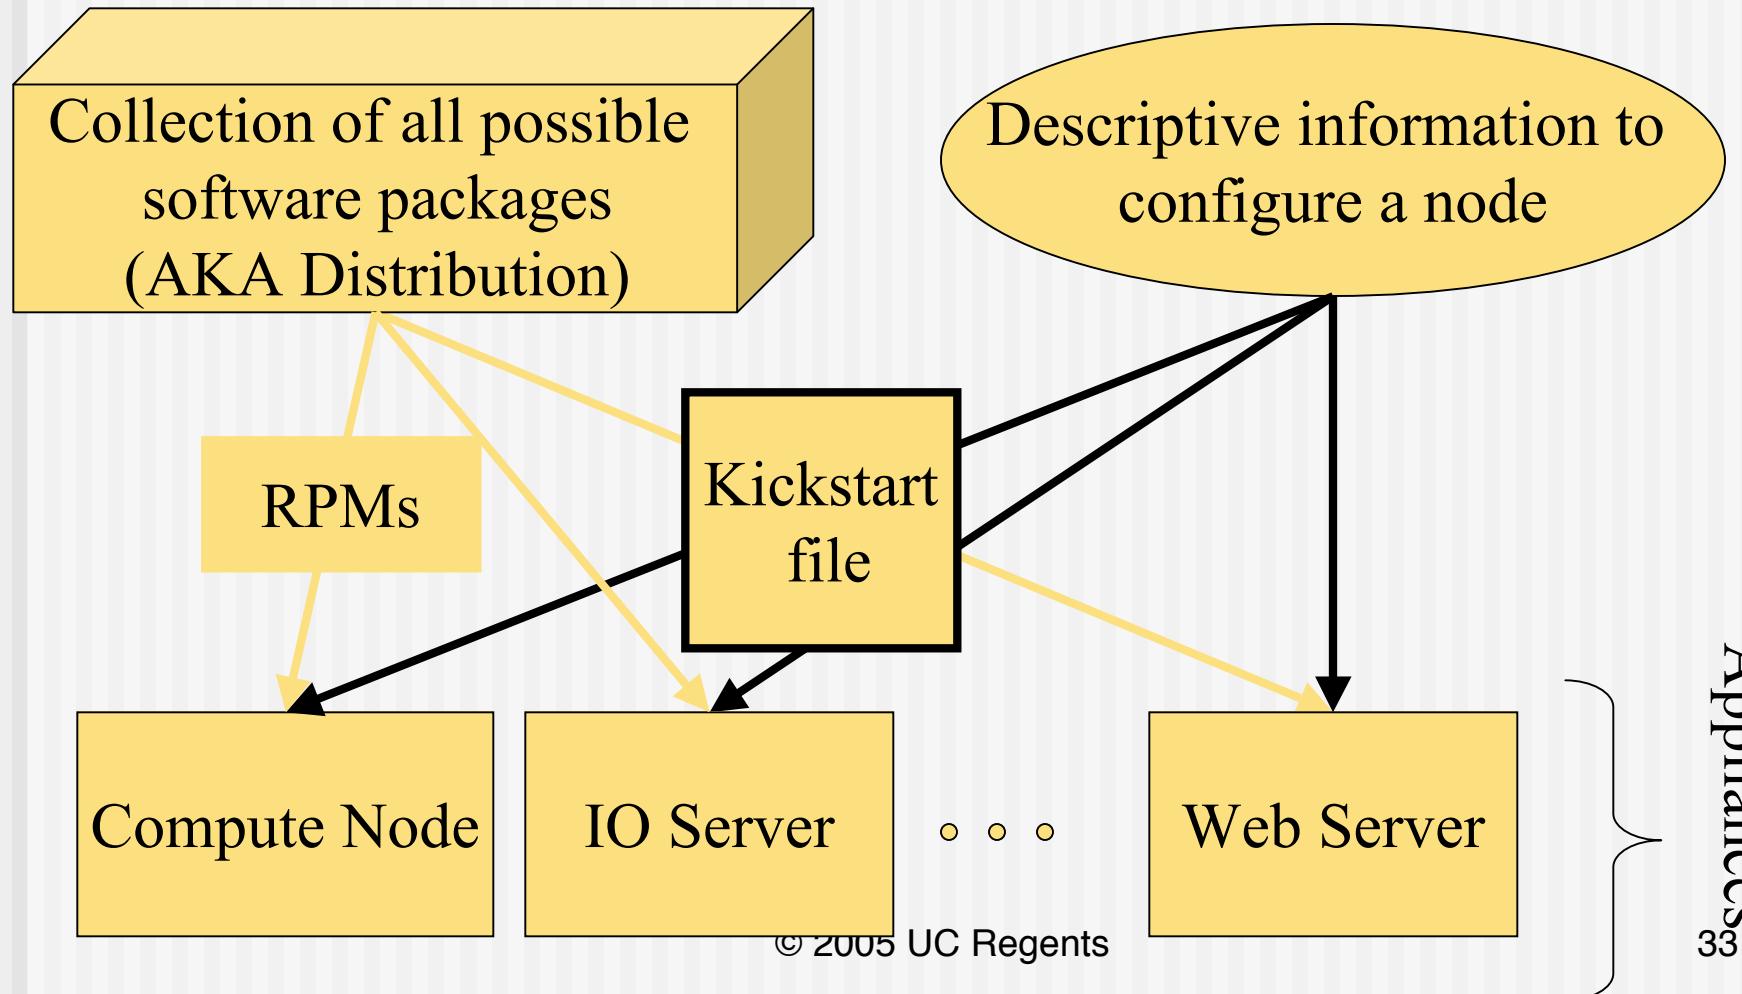
# Ganglia Screenshot



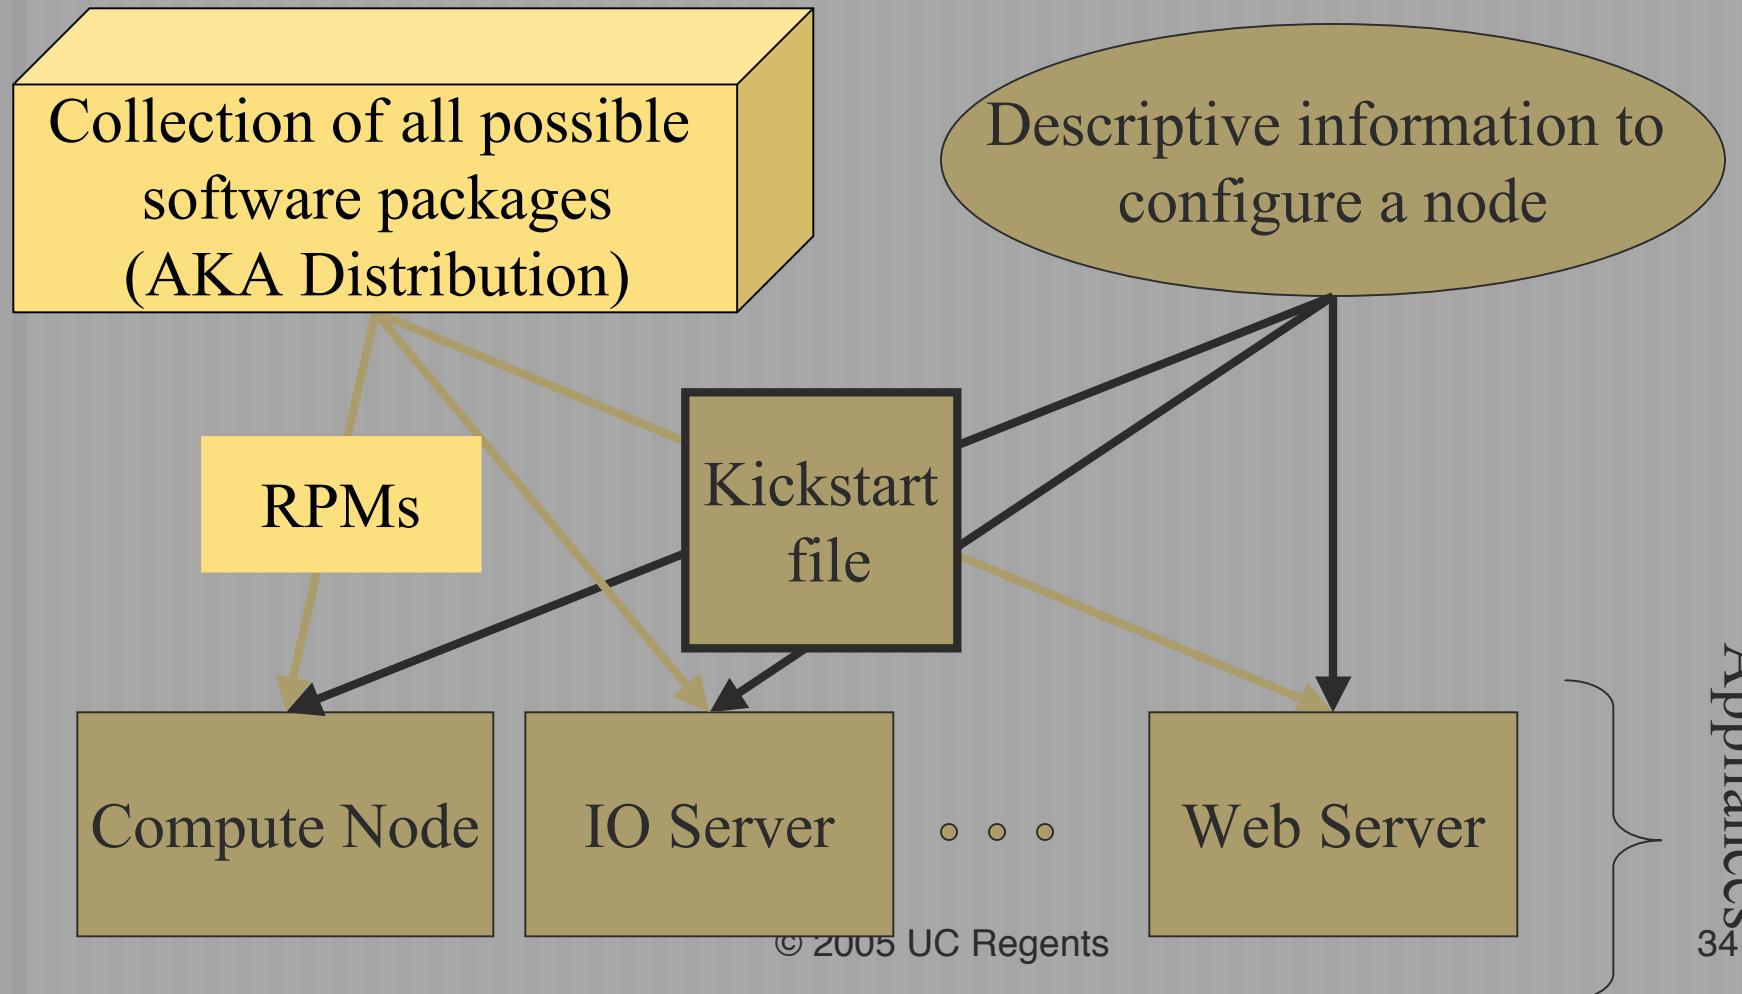
# SCMSWeb Screenshot



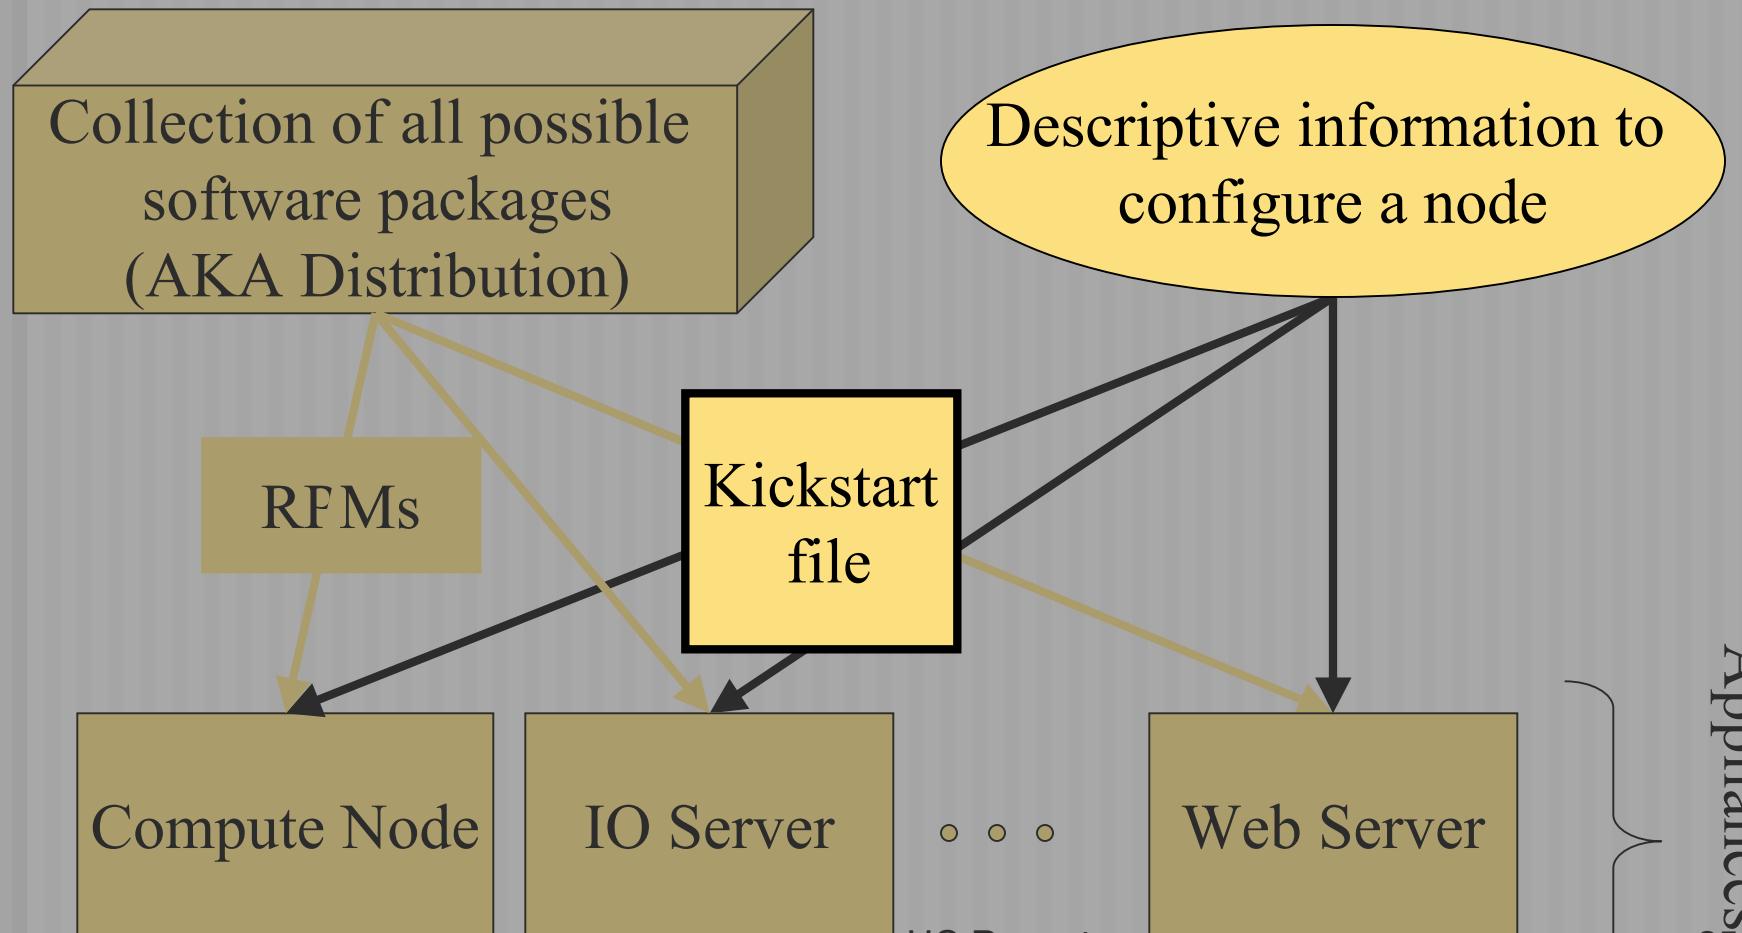
# Software Installation



# Software Repository



# Installation Instructions



# Cluster Software Management

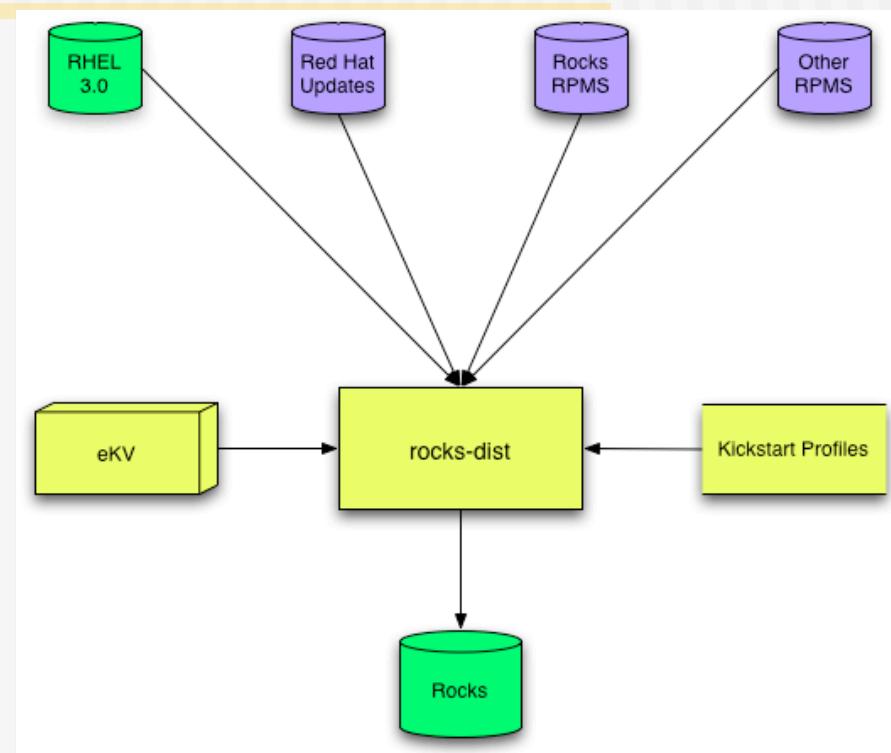
## Software Packages

- ◆ RPMs
  - ➲ Standard Red Hat (desktop) packaged software
  - ➲ Or your own addons
- ◆ Rocks-dist
  - ➲ Manages the RPM repository
  - ➲ This is the distribution

## Software Configuration

- ◆ Tuning RPMs
  - ➲ For clusters
  - ➲ For your site
  - ➲ Other customization
- ◆ XML Kickstart
  - ➲ Programmatic System Building
  - ➲ Scalable

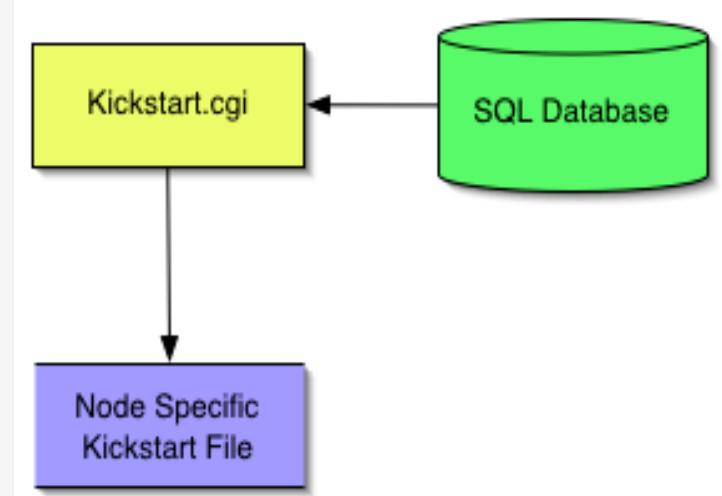
# Building a Rocks Distribution



- ◆ Start with Red Hat
- ◆ Add updates, Rocks (and optional other) software
- ◆ Add Kickstart profiles
- ◆ Modify Red Hat installation boot image
- ◆ Resulting in a Red Hat compatible Rocks distribution

# Kickstart

- ◆ Red Hat's Kickstart
  - ➲ Monolithic flat ASCII file
  - ➲ No macro language
  - ➲ Requires forking based on site information and node type.
- ◆ Rocks XML Kickstart
  - ➲ Decompose a kickstart file into nodes and a graph
    - Graph specifies OO framework
    - Each node specifies a service and its configuration
  - ➲ Macros and SQL for site configuration
  - ➲ Driven from web cgi script



# Kickstart File Sections

- ◆ Main
  - ↳ Disk partitioning
  - ↳ Root password
  - ↳ RPM repository URL
  - ↳ ...
- ◆ Packages
  - ↳ List of RPMs (w/o version numbers)
  - ↳ The repository determines the RPM versions
  - ↳ The kickstart file determines the set of RPMs
- ◆ Pre
  - ↳ Shell scripts run before RPMs are installed
  - ↳ Rarely used (Rocks uses it to enhance kickstart)
- ◆ Post
  - ↳ Shell scripts to cleanup RPM installation
  - ↳ Fixes bugs in packages
  - ↳ Adds local information

# Sample Node File

```
<?xml version="1.0" standalone="no"?>
<!DOCTYPE kickstart SYSTEM "@KICKSTART_DTD@" [<!ENTITY ssh "openssh">]>
<kickstart>
    <description>
        Enable SSH
    </description>

    <package>&ssh;</package>
    <package>&ssh;-clients</package>
    <package>&ssh;-server</package>
    <package>&ssh;-askpass</package>
<post>
    cat &gt; /etc/ssh/ssh_config &lt;&lt; EOF' <!-- default client setup -->
    Host *
        ForwardX11 yes
        ForwardAgent yes
    EOF

    chmod o+rx /root
    mkdir /root/.ssh
    chmod o+rx /root/.ssh

</post>
</kickstart>
```

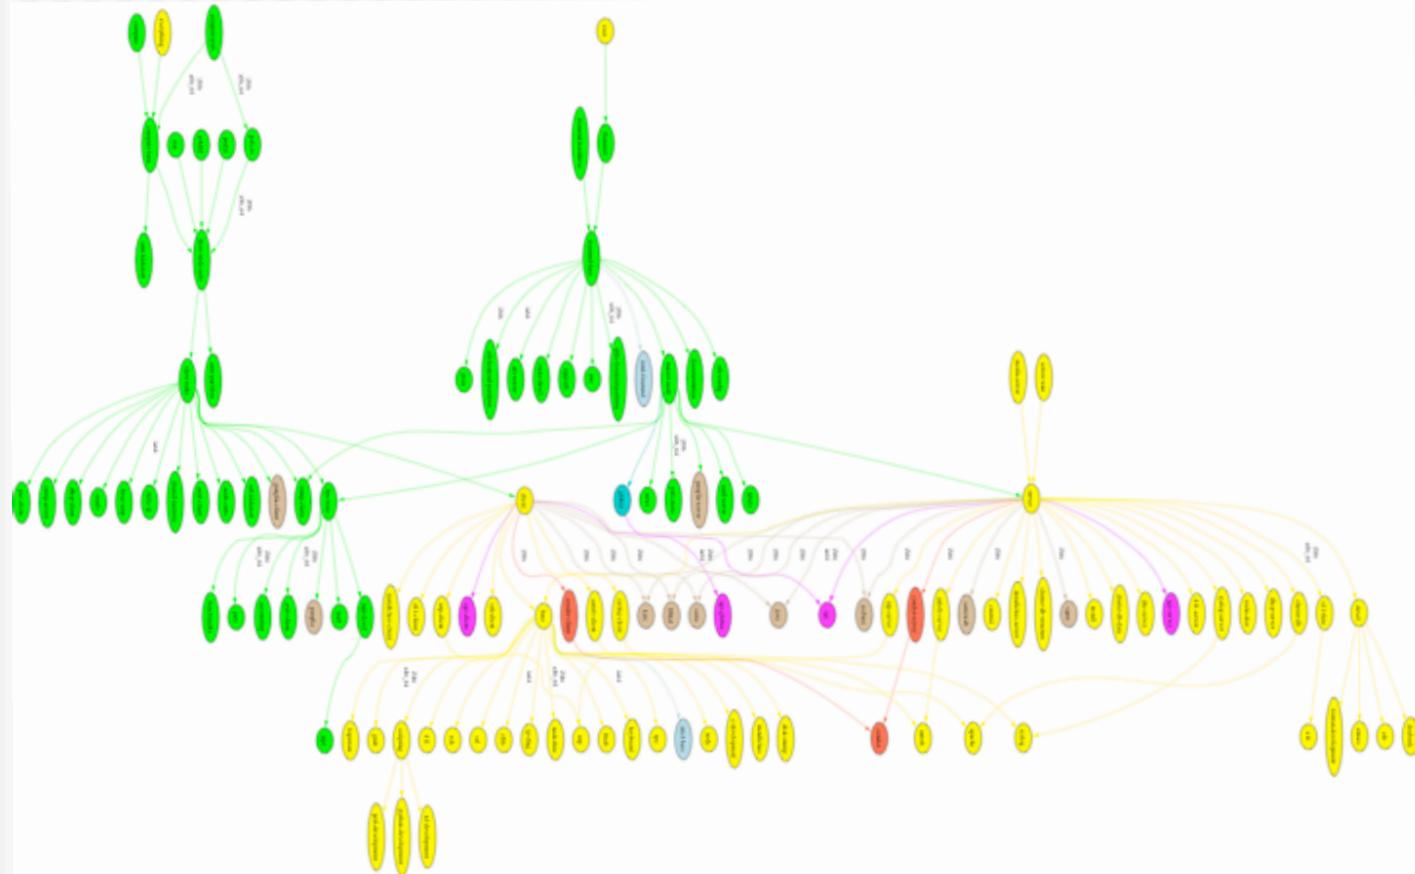
# Sample Graph File

```
<?xml version="1.0" standalone="no"?>

<graph>
    <description>
        Default Graph for NPACI Rocks.
    </description>

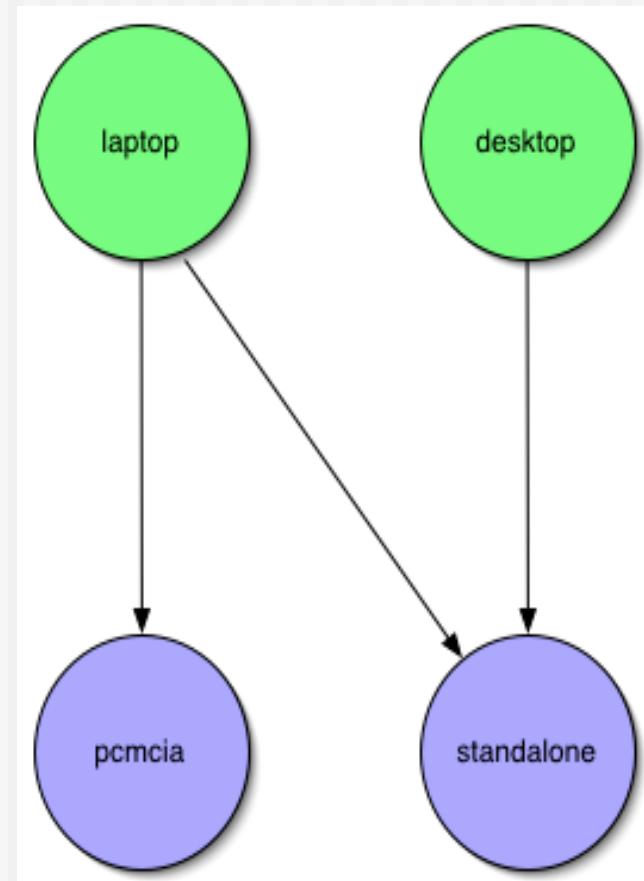
    <edge from="base" to="scripting"/>
    <edge from="base" to="ssh"/>
    <edge from="base" to="ssl"/>
    <edge from="base" to="grub" arch="i386"/>
    <edge from="base" to="elilo" arch="ia64"/>
    ...
    <edge from="node" to="base"/>
    <edge from="node" to="accounting"/>
    <edge from="slave-node" to="node"/>
    <edge from="slave-node" to="nis-client"/>
    <edge from="slave-node" to="autofs-client"/>
    <edge from="slave-node" to="dhcp-client"/>
    <edge from="slave-node" to="snmp-server"/>
    <edge from="slave-node" to="node-certs"/>
    <edge from="compute" to="slave-node"/>
    <edge from="compute" to="usher-server"/>
    <edge from="master-node" to="node"/>
    <edge from="master-node" to="x11"/>
    <edge from="master-node" to="usher-client"/>
</graph>
```

# Kickstart Framework



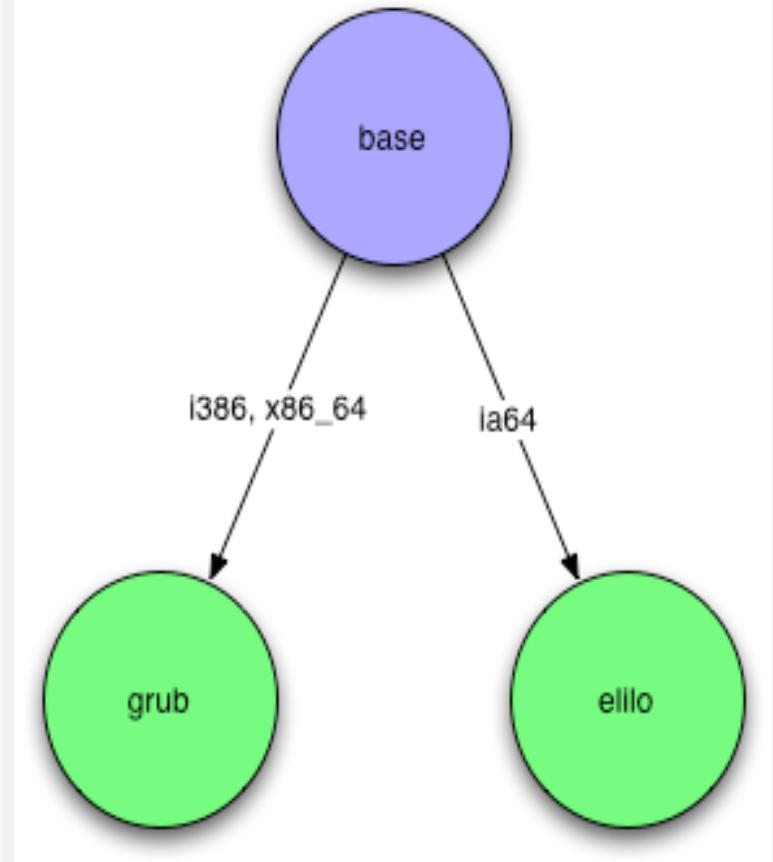
# Appliances

- ◆ Laptop / Desktop
  - ↳ Appliances
  - ↳ Final classes
  - ↳ Node types
- ◆ Desktop IsA
  - ↳ standalone
- ◆ Laptop IsA
  - ↳ standalone
  - ↳ pcmcia
- ◆ Code re-use is good



# Architecture Differences

- ◆ Conditional inheritance
- ◆ Annotate edges with target architectures
- ◆ if i386
  - ⇒ Base IsA grub
- ◆ if ia64
  - ⇒ Base IsA elilo
- ◆ One Graph, Many CPUs
  - ⇒ Heterogeneity is easy
  - ⇒ Not true for SSI or Imaging

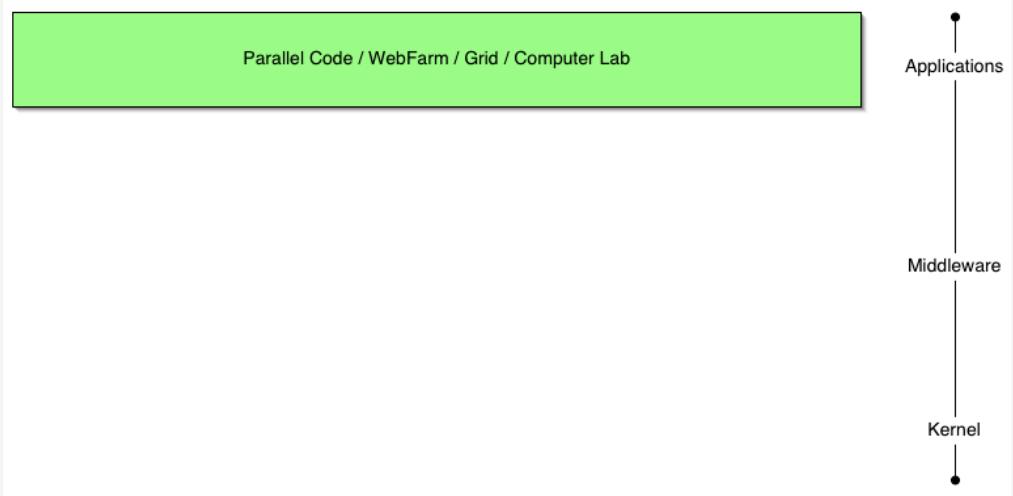


# Optional Drivers

- ◆ PVFS
  - ↪ Parallel Virtual File System
  - ↪ Kernel module built for all nodes
  - ↪ User must decide to enable
- ◆ Myrinet
  - ↪ High Speed and Low Latency Interconnect
  - ↪ GM/MPI for user Applications
  - ↪ Kernel module built for all nodes with Myrinet cards
- ◆ Add your own
  - ↪ Cluster Gigabit Ethernet driver
  - ↪ Infiniband driver

# Application Layer

- ◆ Rocks Rolls
  - ⇒ Optional component
  - ⇒ Created by SDSC
  - ⇒ Created by others
- ◆ Example
  - ⇒ Bio (BLAST)
  - ⇒ Chem (GAMESS)
  - ⇒ Visualization Clusters





**ROCKS**

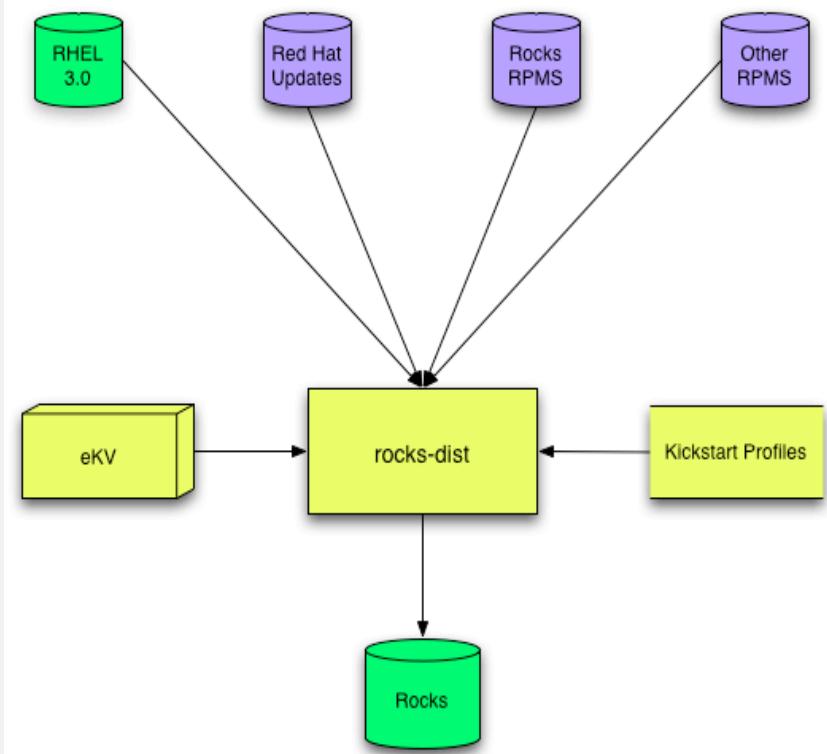
# Building on Top of Rocks

---

## Inheritance and Rolls

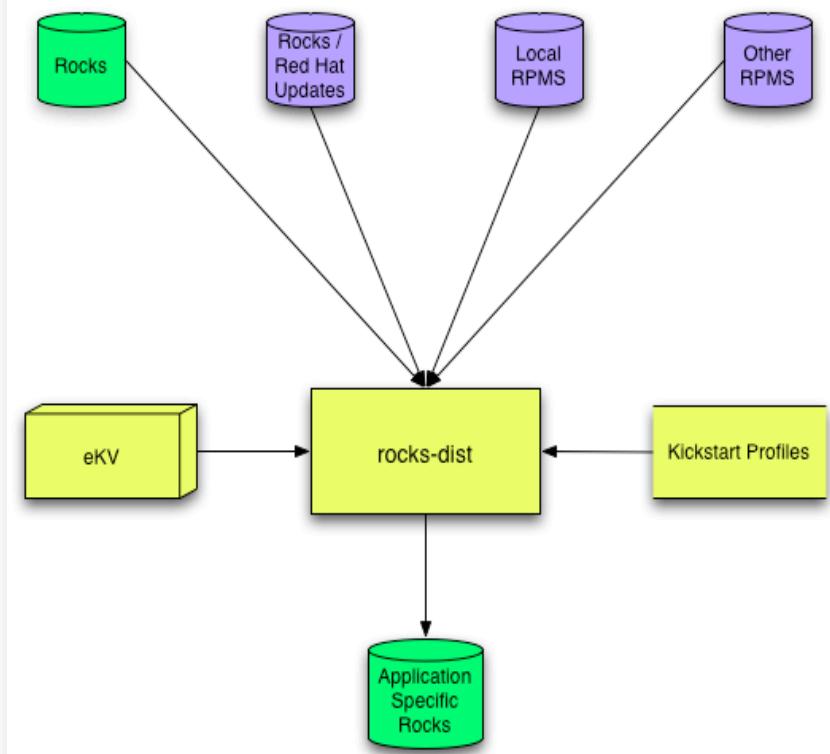
# How Rocks is built

- ◆ Rocks-dist
  - ➲ Merges all RPMs
    - Red Hat
    - Rocks
  - ➲ Resolves versions
  - ➲ Creates Rocks
- ◆ Rocks distribution
  - ➲ Looks just like Red Hat
  - ➲ Cluster optimized Red Hat



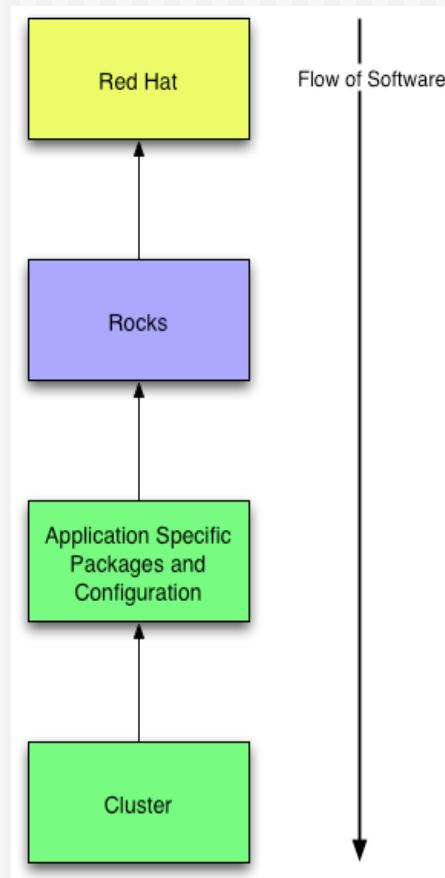
# How You Create Your Own Rocks

- ◆ Rocks-dist
  - ➲ Merges all RPMs
    - Rocks
    - Yours
  - ➲ Resolves versions
  - ➲ Creates Rocks++
- ◆ Your distribution
  - ➲ Looks just like Rocks
  - ➲ Application optimized Rocks



# Extension Through Inheritance

- ◆ UCSD/SDSC Rocks
  - ↳ BIRN
  - ↳ GAMESS Portal
  - ↳ GEON
  - ↳ GriPhyN
- ◆ Commercial
  - ↳ Scalable Systems
  - ↳ Platform Computing
- ◆ Can also override existing functionality
  - ↳ Rocks without NFS?
  - ↳ Rocks for the desktop?



# Rolls

## PICK PACKAGES

- > COMBO #1: PREMIUM
- > COMBO #2: SPORT
- > COMBO #3: COLD WEATHER
- > NEXT STEP



Sport Package (\$1350)

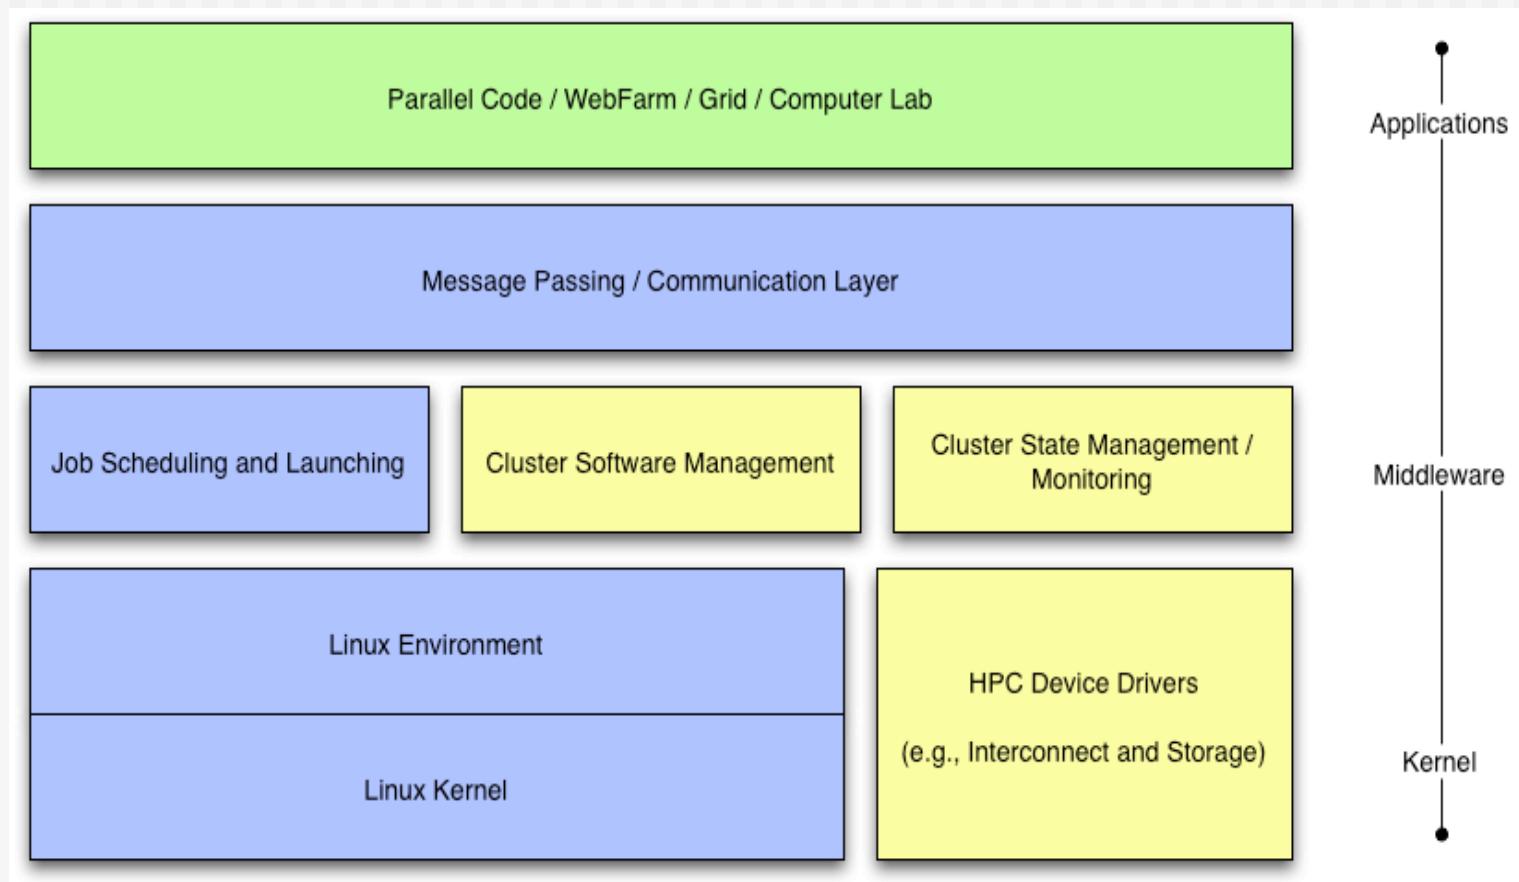
[CLICK IMAGE TO ADD THE SPORT PACKAGE TO YOUR LIST.](#)

THE SPORT PACKAGE WILL ADD:

Dynamic stability control (DSC), bonnet stripes, xenon headlamps with powerwashers, front fog lamps, 17-inch alloy S-lite wheels with 205/45 R17 performance or all-season run-flat tires.

- ◆ Think of a roll as a “package” for a car

# Rolls Break Apart Rocks



# Rocks is What You Make it

- ◆ Motivation
  - ↳ “I’m concerned Rocks is becoming everything for everyone” - rocks mailing list
  - ↳ “Building a cluster should be like ordering a car. I want the sports package, but not the leather seats, …” - z4 owning rocks developer
  - ↳ We need to let go of Rocks but hold onto the core
    - Recruit more external open-source developers
    - Only trust ourselves with fundamental architecture and implementation
  - ↳ We wanted to move the SGE but need to still support PBS
- ◆ Rolls
  - ↳ Optional configuration and software
  - ↳ Just another CD for installed (think application pack)
  - ↳ SGE and PBS are different Rolls
    - User chooses scheduler
    - PBS Roll supported by Norway
    - SGE Roll supported by Singapore (and us)
  - ↳ Rolls give us more flexibility and less work
- ◆ Rocks is done
  - ↳ The core is basically stable and needs continued support
  - ↳ Rolls allow us to develop new ideas
  - ↳ Application Domain specific
- ◆ IEEE Cluster 2004 - “Rolls: Modifying a Standard System Installer to Support User-Customizable Cluster Frontend Appliances”

# Extensible Rocks

- ◆ Over a dozen Rolls already created (e. g.)
  - ↳ SGE, PBS
  - ↳ Grid (NMI stack)
  - ↳ Java
  - ↳ AMD (32bit libraries for Opteron)
  - ↳ Condor
  - ↳ SCE
- ◆ Several third party Rolls have started
  - ↳ Quadrics (rumored)
  - ↳ SGE V6.beta
  - ↳ NIMROD
  - ↳ BIRN
  - ↳ DB2
- ◆ Rocks is done
  - ↳ The core is basically stable and needs continued support
  - ↳ Rolls allow us to develop new ideas
  - ↳ Application Domain specific
  - ↳ For example: Visualization...



**ROCKS**

## Viz Roll

---

Rocks becomes more  
than just compute  
clusters

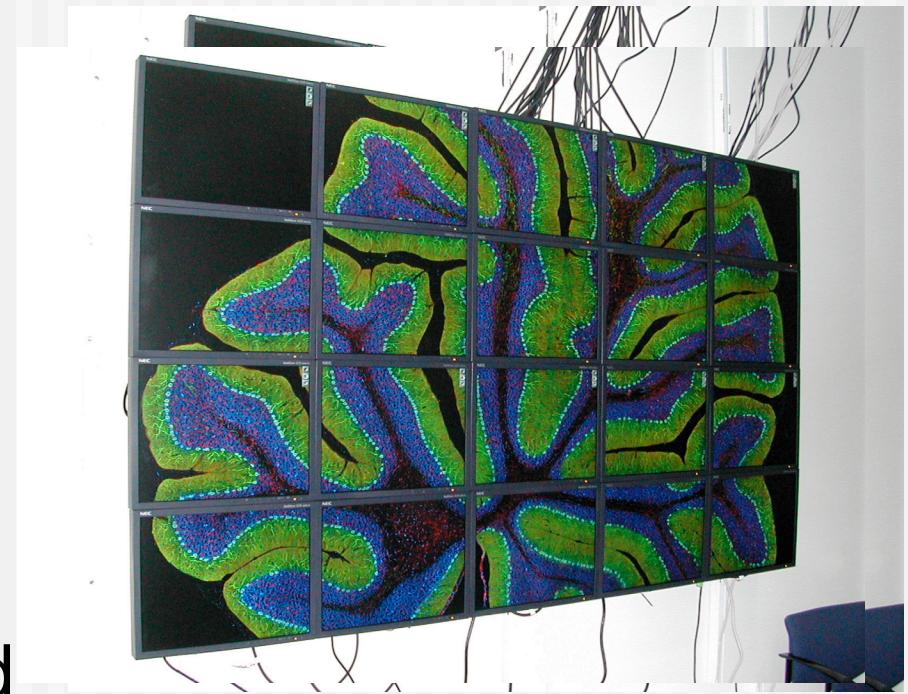
# Early Work: NCSA

- ◆ LCD Cluster
  - ↳ Custom framing
  - ↳ One PC / tile
  - ↳ Portable (luggable)
  - ↳ SC 2001 Demo
- ◆ NCSA Software
  - ↳ Pixel Blaster
  - ↳ Display Wall In-A-Box
  - ↳ OSCAR based
  - ↳ Never fully released



# NCMIR

- ◆ Using Rocks
- ◆ Hand configured a visualization cluster
- ◆ “Administered the machine to the point of instability” - David Lee
- ◆ Automation is needed



# COTS Vis: GeoWall

- ◆ LCD Clusters
  - ➲ One PC / tile
  - ➲ Gigabit Ethernet
  - ➲ Optional Stereo Glasses
  - ➲ Portable
  - ➲ Commercial Frame (Reason)
- ◆ Applications
  - ➲ Large remote sensing
  - ➲ Volume Rendering
  - ➲ Seismic Interpretation
  - ➲ Brain mapping (NCMIR)
- ◆ Electronic Visualization Lab
  - ➲ Jason Leigh (UIC)



# Eye Candy (NCMIR)



# Simple example





**ROCKS**

# Rocks Installation

---

Step by step instruction  
for building your cluster

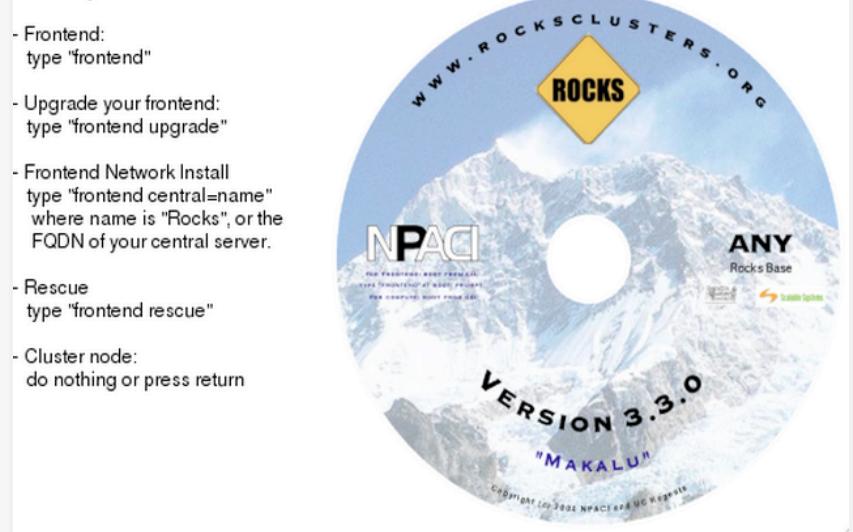
# Frontend Installation

- ◆ Turn on node
- ◆ Insert CDROM
- ◆ Type  
    ⇒ frontend

NPACI Rocks Cluster Distribution

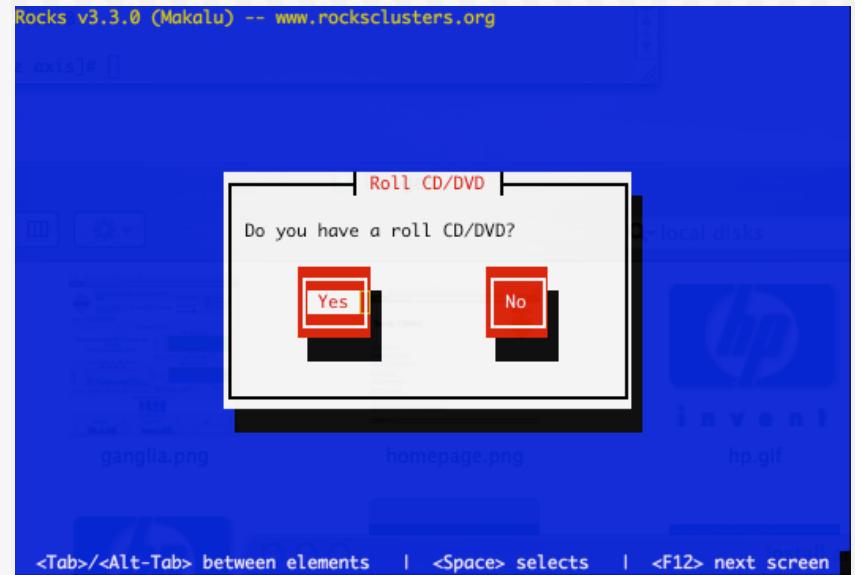
What do you want to kickstart?

- Frontend:  
type "frontend"
- Upgrade your frontend:  
type "frontend upgrade"
- Frontend Network Install  
type "frontend central=name"  
where name is "Rocks", or the  
FQDN of your central server.
- Rescue  
type "frontend rescue"
- Cluster node:  
do nothing or press return



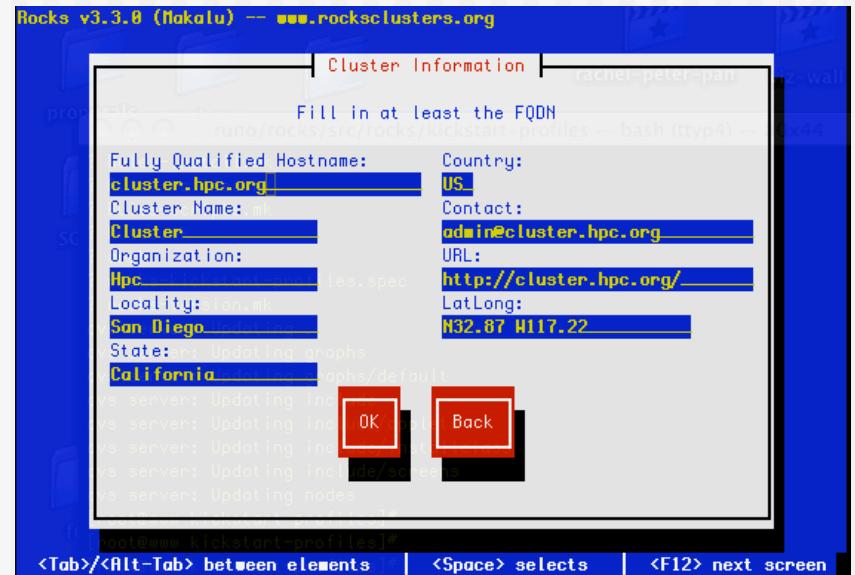
# Rolls

- ◆ Anaconda Starts
- ◆ Asks for Rolls
- ◆ Select “Yes”
- ◆ Insert
  - ⇒ base
  - ⇒ hpc+kernel
  - ⇒ area51+java+grid+sge



# Cluster Information

- ◆ Specific to Rocks
- ◆ Used for Certificates
  - ⇒ SSL/HTTPS
  - ⇒ Globus
- ◆ Hostname
  - ⇒ Must be FQDN
  - ⇒ Must be in DNS
  - ⇒ Must not be an Alias



# Partitioning

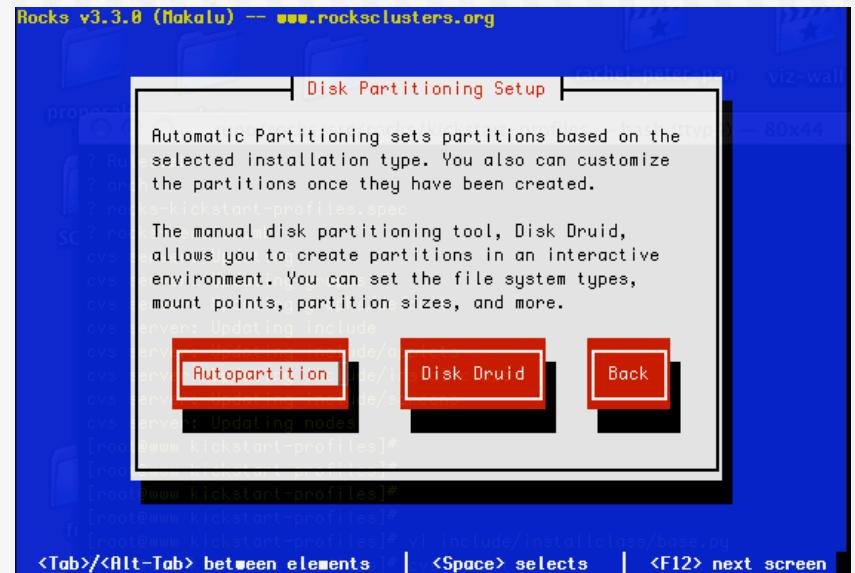
## ◆ Automatic

- ➲ 6GB /
- ➲ 1GB swap
- ➲ Remainder for /export

## ◆ Manual

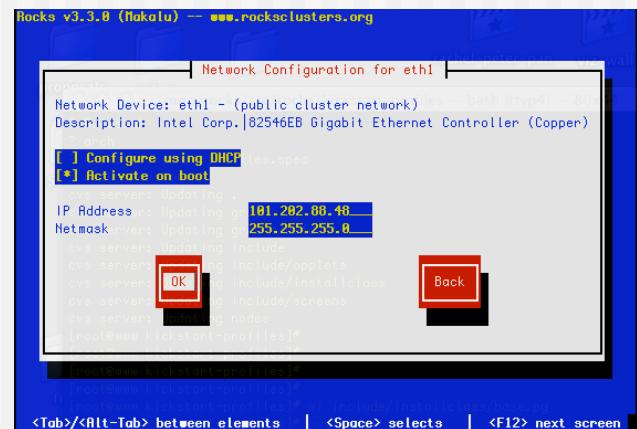
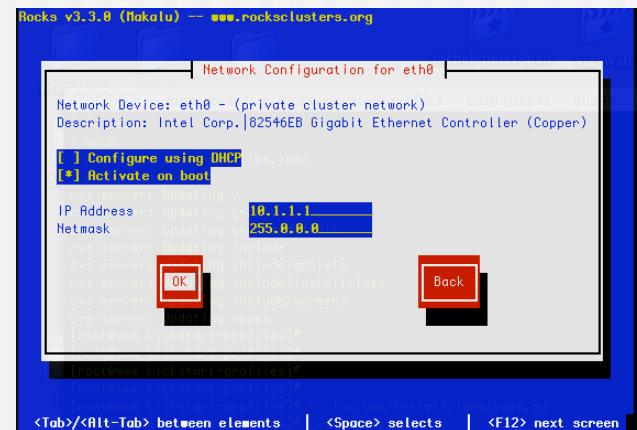
- ➲ You choose
- ➲ Must create a /export

## ◆ Select Wisely



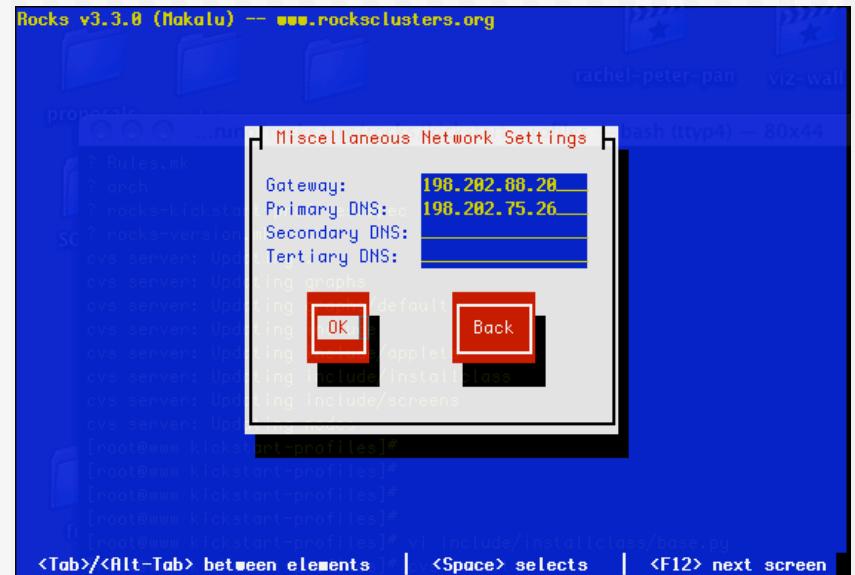
# Networks

- ◆ Private Network
  - ⇒ eth0
  - ⇒ Cluster-side only
- ◆ Public Network
  - ⇒ eth1
  - ⇒ Internet/LAN side
- ◆ You must configure both and have 2 NICs



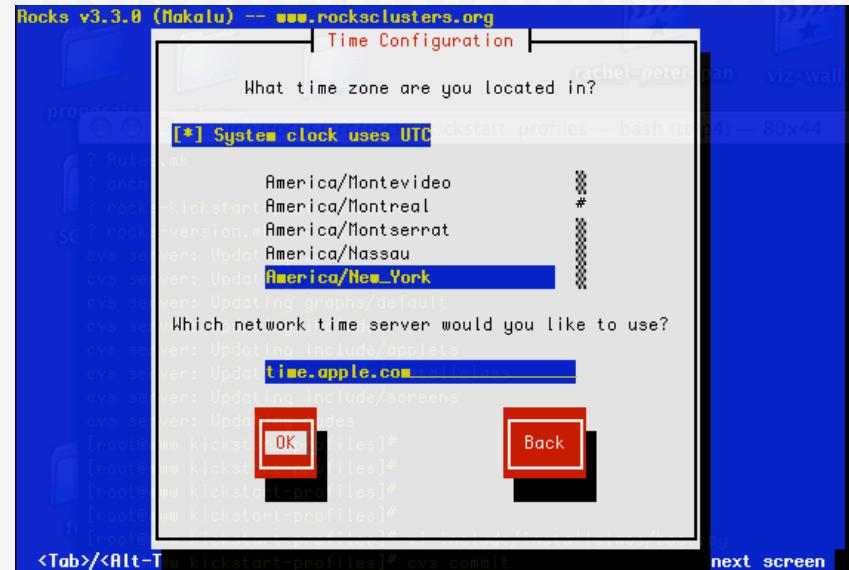
# Gateway

- ◆ Gateway / DNS
  - ↳ Same as any other device on your network
- ◆ All traffic for compute nodes is NATed through the frontend.
- ◆ DNS is only for the frontend, compute nodes use the frontend as their DNS.



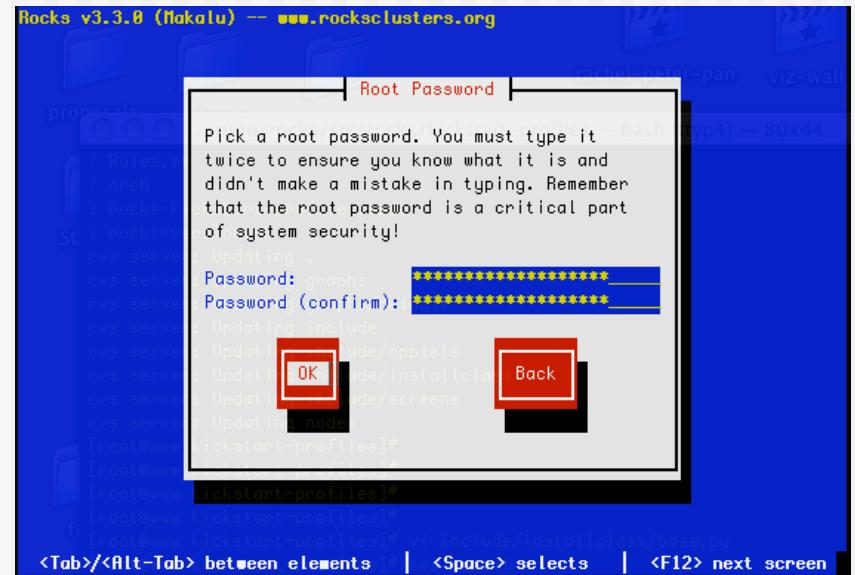
# Network Time Protocol

- ◆ Choose timezone
  - ⇒ UTC is a good choice
  - ⇒ Or localize
- ◆ Default server is
  - ⇒ time.apple.com
  - ⇒ Change it if you wish

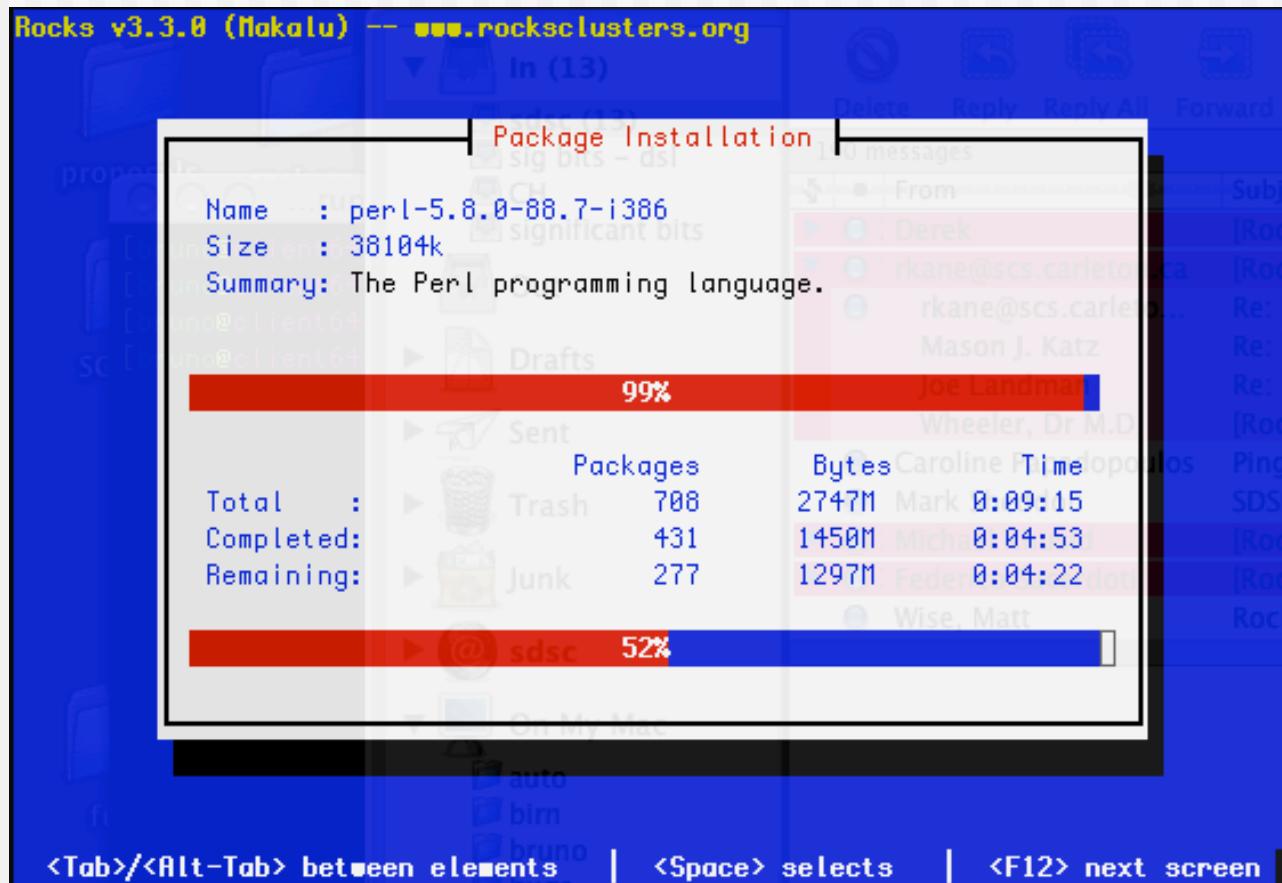


# Root Password

- ◆ Password is secure
  - ➲ Not stored in clear text form anywhere (not in DB)
- ◆ Also used for mysql password

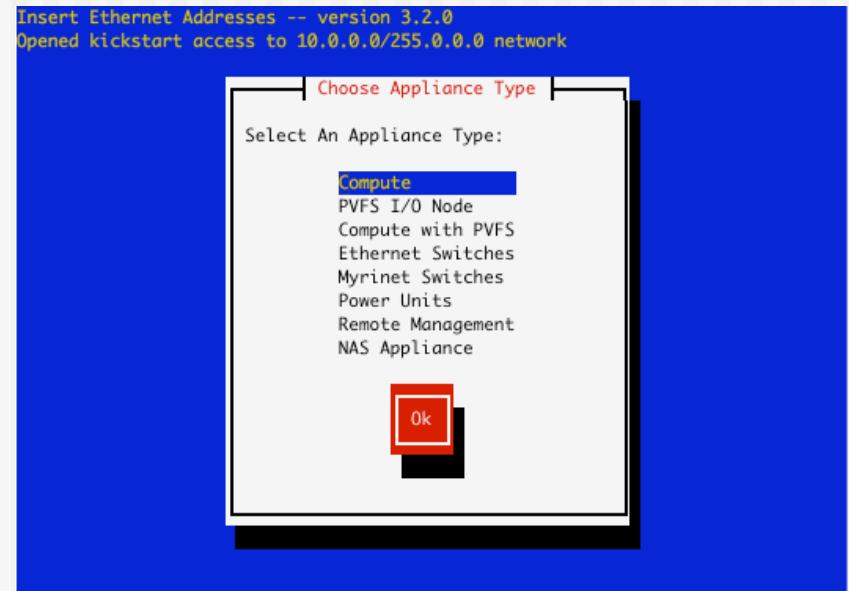


# Installing Packages



# Integrate Compute Nodes

- ◆ Log into Frontend (as root)
- ◆ Run `insert-ethers`
  - ↳ Can choose appliance type
  - ↳ Rolls add new appliance types
  - ↳ For now we will use Compute
- ◆ Turn on first node
  - ↳ Nodes are integrated serially
  - ↳ Need to map machine name to machine location
  - ↳ After we integrate machines can be re-installed in parallel
- ◆ Remote Terminal (ekv)
  - ↳ `ssh compute-0-0 -p2200`



# Discovering Compute-0-0

```
Insert Ethernet Addresses -- version 3.2.0
Opened kickstart access to 10.0.0.0/255.0.0.0 network

[redacted] Inserted Appliances [redacted]
# [redacted]

Press <F10> to quit, press <F11> to force quit
```

```
Insert Ethernet Addresses -- version 3.2.0
Opened kickstart access to 10.0.0.0/255.0.0.0 network

[redacted] Inserted Appliances [redacted]
00:30:c1:a0:ac:25      compute-0-0      ( ) # [redacted]

Press <F10> to quit, press <F11> to force quit
```

```
Insert Ethernet Addresses -- version 3.2.0
Opened kickstart access to 10.0.0.0/255.0.0.0 network

[redacted] Inserted Appliances [redacted]
[redacted] Discovered New Appliance [redacted]
Discovered a new appliance with MAC (00:30:c1:a0:ac:25) [redacted]

Press <F10> to quit, press <F11> to force quit
```

```
[redacted] Inserted Appliances [redacted]
00:30:c1:a0:ac:25      compute-0-0      (*) # [redacted]
```

Retrieved kickstart file

# useradd

```
root@rocks-39:~ — bash (ttyp1)
[root@rocks-39 ~]# useradd mjk
Creating user: mjk
make: Entering directory `/var/411'
/opt/rocks/sbin/411put --comment="#" /etc/auto.home
411 Wrote: /etc/411.d/etc.auto..home
Size: 579/253 bytes (encrypted/plain)
Alert: sent on channel 255.255.255.255:8649 with master 10.1.1.1

/opt/rocks/sbin/411put --comment="#" /etc/passwd
411 Wrote: /etc/411.d/etc.passwd
Size: 2816/1905 bytes (encrypted/plain)
Alert: sent on channel 255.255.255.255:8649 with master 10.1.1.1

/opt/rocks/sbin/411put --comment="#" /etc/shadow
411 Wrote: /etc/411.d/etc.shadow
Size: 1961/1272 bytes (encrypted/plain)
Alert: sent on channel 255.255.255.255:8649 with master 10.1.1.1

/opt/rocks/sbin/411put --comment="#" /etc/group
411 Wrote: /etc/411.d/etc.group
Size: 1236/740 bytes (encrypted/plain)
Alert: sent on channel 255.255.255.255:8649 with master 10.1.1.1

make: Leaving directory `/var/411'
[root@rocks-39 ~]# passwd mjk
Changing password for user mjk.
New UNIX password:
BAD PASSWORD: it is based on a (reversed) dictionary word
Retype new UNIX password:
passwd: all authentication tokens updated successfully.
[root@rocks-39 ~]#
```

# user login

```
mjk@rocks-39:~ — bash (ttyp1)
$~>
$~> ssh concave.rocksclusters.org
mjk@concave.rocksclusters.org's password:
Last login: Mon May 16 19:50:09 2005 from client64-84.sdsc.edu
Rocks Frontend Node - Rocks-39 Cluster
Rocks 4.0.0 (Whitney)
Profile built 13:03 26-Apr-2005

Kickstarted 13:03 26-Apr-2005

It doesn't appear that you have set up your ssh key.
This process will make the files:
  /home/mjk/.ssh/id_rsa.pub
  /home/mjk/.ssh/id_rsa
  /home/mjk/.ssh/authorized_keys

Generating public/private rsa key pair.
Enter file in which to save the key (/home/mjk/.ssh/id_rsa):
Created directory '/home/mjk/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/mjk/.ssh/id_rsa.
Your public key has been saved in /home/mjk/.ssh/id_rsa.pub.
The key fingerprint is:
17:44:24:f3:b7:bd:41:48:4a:82:83:a6:d1:5f:68:af mjk@rocks-39.sdsc.edu
[mjk@rocks-39 ~]$ █
```